

DonorsChoose

DonorsChoose.org receives hundreds of thousands of project proposals each year for classroom projects in need of funding. Right now, a large number of volunteers is needed to manually screen each submission before it's approved to be posted on the DonorsChoose.org website.

Next year, DonorsChoose.org expects to receive close to 500,000 project proposals. As a result, there are three main problems they need to solve:

- How to scale current manual processes and resources to screen 500,000 projects so that they can be posted as quickly and as efficiently as possible
- How to increase the consistency of project vetting across different volunteers to improve the experience for teachers
- How to focus volunteer time on the applications that need the most assistance

The goal of the competition is to predict whether or not a DonorsChoose.org project proposal submitted by a teacher will be approved, using the text of project descriptions as well as additional metadata about the project, teacher, and school. DonorsChoose.org can then use this information to identify projects most likely to need further review before approval.

About the DonorsChoose Data Set

The `train.csv` data set provided by DonorsChoose contains the following features:

Feature	Description
<code>project_id</code>	A unique identifier for the proposed project. Example: p036502
<code>project_title</code>	Title of the project. Examples: <code>Art Will Make You Happy!</code> <code>First Grade Fun</code>
<code>project_grade_category</code>	Grade level of students for which the project is targeted. One of the following enumerated values: <code>Grades PreK-2</code> <code>Grades 3-5</code> <code>Grades 6-8</code> <code>Grades 9-12</code>
<code>project_subject_categories</code>	One or more (comma-separated) subject categories for the project from the following enumerated list of values: <code>Applied Learning</code> <code>Care & Hunger</code> <code>Health & Sports</code> <code>History & Civics</code> <code>Literacy & Language</code> <code>Math & Science</code> <code>Music & The Arts</code> <code>Special Needs</code> <code>Warmth</code> Examples: <code>Music & The Arts</code> <code>Literacy & Language, Math & Science</code>
<code>school_state</code>	State where school is located (Two-letter U.S. postal code). Example: WY
<code>project_subject_subcategories</code>	One or more (comma-separated) subject subcategories for the project. Examples: <code>Literacy</code> <code>Literature & Writing, Social Sciences</code>
<code>project_resource_summary</code>	An explanation of the resources needed for the project. Example: <code>My students need hands on literacy materials to manage sensory needs!</code>
<code>project_essay_1</code>	First application essay*
<code>project_essay_2</code>	Second application essay*
<code>project_essay_3</code>	Third application essay*

Feature	Description
project_essay_4	Fourth application essay
project_submitted_datetime	Datetime when project application was submitted. Example: 2016-04-28 12:43:56.245
teacher_id	A unique identifier for the teacher of the proposed project. Example: bdf8baa8fedef6bfeec7ae4ff1c15c56
teacher_prefix	Teacher's title. One of the following enumerated values: <ul style="list-style-type: none"> nan Dr. Mr. Mrs. Ms. Teacher.
teacher_number_of_previously_posted_projects	Number of project applications previously submitted by the same teacher. Example: 2

* See the section **Notes on the Essay Data** for more details about these features.

Additionally, the `resources.csv` data set provides more data about the resources required for each project. Each line in this file represents a resource required by a project:

Feature	Description
id	A <code>project_id</code> value from the <code>train.csv</code> file. Example: p036502
description	Description of the resource. Example: Tenor Saxophone Reeds, Box of 25
quantity	Quantity of the resource required. Example: 3
price	Price of the resource required. Example: 9.95

Note: Many projects require multiple resources. The `id` value corresponds to a `project_id` in `train.csv`, so you use it as a key to retrieve all resources needed for a project:

The data set contains the following label (the value you will attempt to predict):

Label	Description
project_is_approved	A binary flag indicating whether DonorsChoose approved the project. A value of 0 indicates the project was not approved, and a value of 1 indicates the project was approved.

Notes on the Essay Data

Prior to May 17, 2016, the prompts for the essays were as follows:

- `__project_essay_1__`: "Introduce us to your classroom"
- `__project_essay_2__`: "Tell us more about your students"
- `__project_essay_3__`: "Describe how your students will use the materials you're requesting"
- `__project_essay_3__`: "Close by sharing why your project will make a difference"

Starting on May 17, 2016, the number of essays was reduced from 4 to 2, and the prompts for the first 2 essays were changed to the following:

- `__project_essay_1__`: "Describe your students: What makes your students special? Specific details about their background, your neighborhood, and your school are all helpful."
- `__project_essay_2__`: "About your project: How will these materials make a difference in your students' learning and improve their school lives?"

For all projects with `project_submitted_datetime` of 2016-05-17 and later, the values of `project_essay_3` and `project_essay_4` will be NaN.

In [0]:

```
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
warnings.simplefilter("ignore")
warnings.warn("ignore")
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
```

```

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn import model_selection
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

```

In [0]:

```

from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve, auc
from sklearn import preprocessing

```

In [4]:

```

from keras.preprocessing.text import one_hot
from keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers.core import Activation, Dropout, Dense, Reshape
from keras.layers import Flatten, LSTM, Lambda
from keras.models import Model
from keras.layers.embeddings import Embedding
from keras.preprocessing.text import Tokenizer
from keras.layers import Input
from keras.layers import Concatenate
from keras.utils import to_categorical
from keras.layers import Conv1D, MaxPooling1D

```

Using TensorFlow backend.

1.1 Reading Data

In [0]:

```

#https://stackabuse.com/python-for-nlp-creating-multi-data-type-classification-models-with-keras/
#https://www.pyimagesearch.com/2019/01/21/regression-with-keras/
#https://github.com/mmortazavi/EntityEmbedding-Working_Example/blob/master/EntityEmbedding.ipynb
#https://www.pyimagesearch.com/2019/02/04/keras-multiple-inputs-and-mixed-data/
#https://machinelearningmastery.com/cnn-models-for-human-activity-recognition-time-series-classification/

```

In [0]:

```

preprocessed_data = pd.read_csv('preprocessed_data.csv')

```

In [6]:

```

print("Number of data points in preprocessed data", preprocessed_data.shape)

```

Number of data points in preprocessed data: (100248, 8)

number of data points in preprocessed data (109246, 9)

In [7]:

```
preprocessed_data=preprocessed_data.sample(n=100000)
preprocessed_data.head()
```

Out[7]:

school_state	teacher_prefix	project_grade_category	teacher_number_of_previously_posted_projects	project_is_approved	clean	
6876	ny	ms	grades_prek_2	16	1	n
105702	wv	ms	grades_9_12	6	1	h
62830	ca	mrs	grades_3_5	0	1	litera
46841	tn	ms	grades_prek_2	0	1	litera n
30042	wi	mr	grades_6_8	3	1	

In [0]:

```
X=preprocessed_data.drop(columns=['project_is_approved'],axis=1)
y=preprocessed_data['project_is_approved']
```

In [0]:

```
label_encoder = preprocessing.LabelEncoder()
y = label_encoder.fit_transform(y)
```

In [0]:

```
X_1, X_test, y_1, y_test = model_selection.train_test_split(X, y, test_size=0.2, random_state=0, stratify=y)

# split the train data set into cross validation train and cross validation test
X_train, X_cv, y_train, y_cv = model_selection.train_test_split(X_1, y_1, test_size=0.2, random_state=0, stratify=y_1)
```

In [0]:

```
y_train = to_categorical(y_train)
y_cv = to_categorical(y_cv)
y_test = to_categorical(y_test)
```

In [0]:

```
tokenizer = Tokenizer()
```

```
tokenizer.fit_on_texts(X_train['essay'].values)

X1_tr = np.array(tokenizer.texts_to_sequences(X_train['essay'].values))
X1_cv = np.array(tokenizer.texts_to_sequences(X_cv['essay'].values))
X1_test = np.array(tokenizer.texts_to_sequences(X_test['essay'].values))
```

In [0]:

```
vocab_size = len(tokenizer.word_index) + 1

maxlen = 200

X1_tr = pad_sequences(X1_tr, padding='post', maxlen=maxlen)
X1_cv = pad_sequences(X1_cv, padding='post', maxlen=maxlen)
X1_test = pad_sequences(X1_test, padding='post', maxlen=maxlen)
```

In [14]:

```
print(X1_tr.shape)
print(X1_cv.shape)
print(X1_test.shape)
```

```
(64000, 200)
(16000, 200)
(20000, 200)
```

In [0]:

```
with open('glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())
```

In [0]:

```
embeddings_dictionary = dict()
for word in glove_words:
    vector_dimensions = model[word]
    embeddings_dictionary[word] = vector_dimensions
```

In [0]:

```
embedding_matrix = np.zeros((vocab_size, 300))
for word, index in tokenizer.word_index.items():
    embedding_vector = embeddings_dictionary.get(word)
    if embedding_vector is not None:
        embedding_matrix[index] = embedding_vector
```

In [18]:

```
embedding_matrix.shape
```

Out[18]:

```
(45666, 300)
```

In [19]:

```
input_1 = Input(shape=(maxlen,), name='essay_input')
print(input_1.shape)
input_1_embedding = Embedding(vocab_size, 300, weights=[embedding_matrix], trainable=False)(input_1)
print(input_1_embedding.shape)
input_1_lstm = LSTM(128, return_sequences=True)(input_1_embedding)
print(input_1_lstm.shape)
input_1_flatten = Flatten()(input_1_lstm)
print(input_1_flatten.shape)
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please us

```
packages/keras/backend/tensorflow_backend.py:191: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.
```

```
(?, 200)
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:190: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:203: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:207: The name tf.global_variables is deprecated. Please use tf.compat.v1.global_variables instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:216: The name tf.is_variable_initialized is deprecated. Please use tf.compat.v1.is_variable_initialized instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:223: The name tf.variables_initializer is deprecated. Please use tf.compat.v1.variables_initializer instead.
```

```
(?, 200, 300)
(?, ?, 128)
(?, ?)
```

In [0]:

```
categoricals=['school_state','teacher_prefix','project_grade_category','clean_categories','clean_subcategories']
numericals=['teacher_number_of_previously_posted_projects','price']
```

In [21]:

```
embed_cols=[i for i in X_train[categoricals]]

for i in embed_cols:
    print(i,X_train[i].nunique())
```

```
school_state 51
teacher_prefix 5
project_grade_category 4
clean_categories 51
clean_subcategories 384
```

In [0]:

```
from sklearn.feature_extraction.text import CountVectorizer
```

In [23]:

```
vectorizer = CountVectorizer()
X2_tr = vectorizer.fit_transform(X_train['school_state'].values).toarray()
X2_cv = vectorizer.transform(X_cv['school_state'].values).toarray()
X2_test = vectorizer.transform(X_test['school_state'].values).toarray()
cat_emb_name= 'school_state_Embedding'
no_of_unique_cat = X_train['school_state'].nunique()
embedding_size = int(min(np.ceil((no_of_unique_cat)/2), 50 ))
input_2 = Input(shape=(X2_tr.shape[1]),name='school_state_input')
```

```

print(input_2.shape)
input_2_embedding = Embedding(no_of_unique_cat, embedding_size, input_length=X2_tr.shape[1],
name=cat_emb_name)(input_2)
print(input_2_embedding.shape)
input_2_flatten=Flatten()(input_2_embedding)
print(input_2_flatten.shape)

```

```

(?, 51)
(?, 51, 26)
(?, ?)

```

In [24]:

```

print(X2_tr.shape)
print(X2_cv.shape)
print(X2_test.shape)

```

```

(64000, 51)
(16000, 51)
(20000, 51)

```

In [25]:

```

vectorizer = CountVectorizer()
X3_tr = vectorizer.fit_transform(X_train['teacher_prefix'].values.astype('U')).toarray()
X3_cv = vectorizer.transform(X_cv['teacher_prefix'].values.astype('U')).toarray()
X3_test = vectorizer.transform(X_test['teacher_prefix'].values.astype('U')).toarray()
cat_emb_name= 'teacher_prefix_Embedding'
no_of_unique_cat = X_train['teacher_prefix'].nunique()
embedding_size = int(min(np.ceil((no_of_unique_cat)/2), 50 ))
input_3 = Input(shape=(X3_tr.shape[1],), name='teacher_prefix_input')
print(input_3.shape)
input_3_embedding = Embedding(no_of_unique_cat, embedding_size, input_length=X3_tr.shape[1],
name=cat_emb_name)(input_3)
print(input_3_embedding.shape)
input_3_flatten=Flatten()(input_3_embedding)
print(input_3_flatten.shape)

```

```

(?, 5)
(?, 5, 3)
(?, ?)

```

In [25]:

```

print(X3_tr.shape)
print(X3_cv.shape)
print(X3_test.shape)

```

```

(64000, 5)
(16000, 5)
(20000, 5)

```

In [26]:

```

vectorizer = CountVectorizer()
X4_tr = vectorizer.fit_transform(X_train['project_grade_category'].values).toarray()
X4_cv = vectorizer.transform(X_cv['project_grade_category'].values).toarray()
X4_test = vectorizer.transform(X_test['project_grade_category'].values).toarray()
cat_emb_name= 'project_grade_category_Embedding'
no_of_unique_cat = X_train['project_grade_category'].nunique()
embedding_size = int(min(np.ceil((no_of_unique_cat)/2), 50 ))
input_4 = Input(shape=(X4_tr.shape[1],), name='project_grade_category_input')
print(input_4.shape)
input_4_embedding = Embedding(no_of_unique_cat,
embedding_size, input_length=X4_tr.shape[1], name=cat_emb_name)(input_4)
print(input_4_embedding.shape)
input_4_flatten=Flatten()(input_4_embedding)
print(input_4_flatten.shape)

```

```

(?, 4)

```

```
..., 2)
(?, 4, 2)
(?, ?)
```

In [27]:

```
print(X4_tr.shape)
print(X4_cv.shape)
print(X4_test.shape)
```

```
(64000, 4)
(16000, 4)
(20000, 4)
```

In [27]:

```
vectorizer = CountVectorizer()
X5_tr = vectorizer.fit_transform(X_train['clean_categories'].values).toarray()
X5_cv = vectorizer.transform(X_cv['clean_categories'].values).toarray()
X5_test = vectorizer.transform(X_test['clean_categories'].values).toarray()
cat_emb_name= 'clean_categories_Embedding'
no_of_unique_cat = X_train['clean_categories'].nunique()
embedding_size = int(min(np.ceil((no_of_unique_cat)/2), 50 ))
input_5 = Input(shape=(X5_tr.shape[1],), name='clean_categories_input')
print(input_5.shape)
input_5_embedding = Embedding(no_of_unique_cat, embedding_size, input_length=X5_tr.shape[1],
name=cat_emb_name)(input_5)
print(input_5_embedding.shape)
input_5_flatten=Flatten()(input_5_embedding)
print(input_5_flatten.shape)
```

```
(?, 9)
(?, 9, 26)
(?, ?)
```

In [29]:

```
print(X5_tr.shape)
print(X5_cv.shape)
print(X5_test.shape)
```

```
(64000, 9)
(16000, 9)
(20000, 9)
```

In [28]:

```
vectorizer = CountVectorizer()
X6_tr = vectorizer.fit_transform(X_train['clean_subcategories'].values).toarray()
X6_cv = vectorizer.transform(X_cv['clean_subcategories'].values).toarray()
X6_test = vectorizer.transform(X_test['clean_subcategories'].values).toarray()
cat_emb_name= 'clean_subcategories_Embedding'
no_of_unique_cat = X_train['clean_subcategories'].nunique()
embedding_size = int(min(np.ceil((no_of_unique_cat)/2), 50 ))
input_6 = Input(shape=(X6_tr.shape[1],), name='clean_subcategories_input')
print(input_6.shape)
input_6_embedding = Embedding(no_of_unique_cat, embedding_size, input_length=X6_tr.shape[1],
name=cat_emb_name)(input_6)
print(input_6_embedding.shape)
input_6_flatten=Flatten()(input_6_embedding)
print(input_6_flatten.shape)
```

```
(?, 30)
(?, 30, 50)
(?, ?)
```

In [31]:

```
print(X6_tr.shape)
print(X6_cv.shape)
```



```
print(X6_test.shape)
```

```
(64000, 30)
(16000, 30)
(20000, 30)
```

In [29]:

```
X7_tr = preprocessing.normalize(X_train[['teacher_number_of_previously_posted_projects', 'price']]
)
X7_cv = preprocessing.normalize(X_cv[['teacher_number_of_previously_posted_projects', 'price']]
X7_test = preprocessing.normalize(X_test[['teacher_number_of_previously_posted_projects', 'price']]
)
input_7 = Input(shape=(len(X_train[numericals].columns),), name='numerical_input')
print(input_7.shape)
input_7_dense = Dense(128)(input_7)
print(input_7_dense.shape)
```

```
(?, 2)
(?, 128)
```

In [33]:

```
print(X7_tr.shape)
print(X7_cv.shape)
print(X7_test.shape)
```

```
(64000, 2)
(16000, 2)
(20000, 2)
```

In [0]:

```
from keras.regularizers import l2
#At the end we concatenate altogether and add other Dense layers
output_1 = Concatenate()
([input_1_flatten, input_2_flatten, input_3_flatten, input_4_flatten, input_5_flatten, input_6_flatten,
input_7_dense])
output_1 = Dense(256, activation='relu', kernel_initializer='he_normal')(output_1)
output_1= Dropout(0.2)(output_1)
output_1 = Dense(128, activation='relu', kernel_initializer='he_normal')(output_1)
output_1= Dropout(0.2)(output_1)
output_1 = Dense(64, activation='relu', kernel_initializer='he_normal')(output_1)
output_1= Dropout(0.3)(output_1)
output_1 = Dense(2, activation='softmax')(output_1)
```

In [0]:

```
#https://stackoverflow.com/questions/41032551/how-to-compute-receiving-operating-characteristic-ro
c-and-auc-in-keras
import tensorflow as tf
def auROC(y_true, y_pred):
    return tf.py_func(roc_auc_score, (y_true, y_pred), tf.double)
```

In [142]:

```
from keras import optimizers
model = Model(inputs=[input_1, input_2, input_3, input_4, input_5, input_6, input_7], outputs=output_1)
model.compile(loss='binary_crossentropy', optimizer=optimizers.SGD(lr=0.01, decay=1e-3, momentum=0.9
), metrics=[auROC])
model.summary()
```

Model: "model_21"


Layer (type)	Output Shape	Param #	Connected to
essay_input (InputLayer)	(None, 200)	0	
embedding_1 (Embedding)	(None, 200, 300)	13699800	essay_input[0][0]

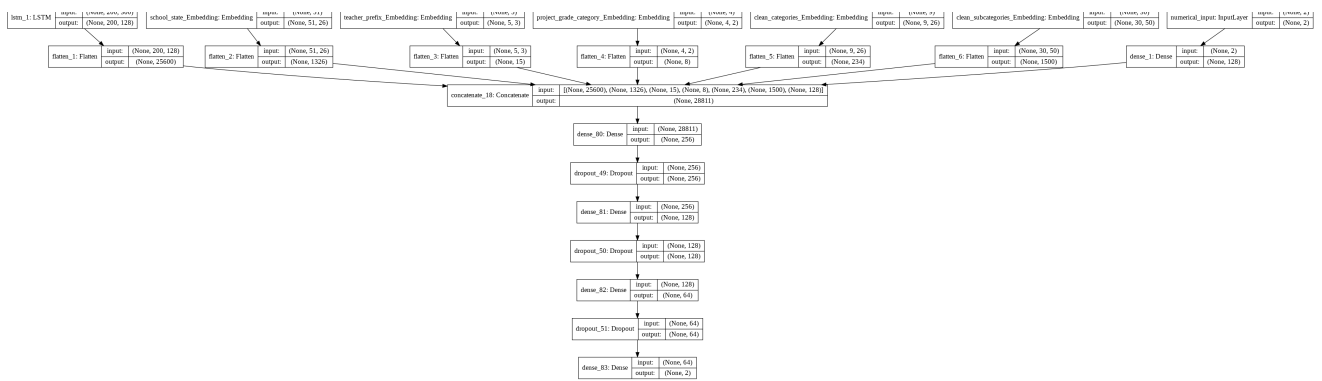
school_state_input (InputLayer)	(None, 51)	0	
teacher_prefix_input (InputLayer)	(None, 5)	0	
project_grade_category_input (InputLayer)	(None, 4)	0	
clean_categories_input (InputLayer)	(None, 9)	0	
clean_subcategories_input (InputLayer)	(None, 30)	0	
lstm_1 (LSTM)	(None, 200, 128)	219648	embedding_1[0][0]
school_state_Embedding (Embedding)	(None, 51, 26)	1326	school_state_input[0][0]
teacher_prefix_Embedding (Embedding)	(None, 5, 3)	15	teacher_prefix_input[0][0]
project_grade_category_Embedding (Embedding)	(None, 4, 2)	8	project_grade_category_input[0][0]
clean_categories_Embedding (Embedding)	(None, 9, 26)	1326	clean_categories_input[0][0]
clean_subcategories_Embedding (Embedding)	(None, 30, 50)	19200	clean_subcategories_input[0][0]
numerical_input (InputLayer)	(None, 2)	0	
flatten_1 (Flatten)	(None, 25600)	0	lstm_1[0][0]
flatten_2 (Flatten)	(None, 1326)	0	school_state_Embedding[0][0]
flatten_3 (Flatten)	(None, 15)	0	teacher_prefix_Embedding[0][0]
flatten_4 (Flatten)	(None, 8)	0	project_grade_category_Embedding[0][0]
flatten_5 (Flatten)	(None, 234)	0	clean_categories_Embedding[0][0]
flatten_6 (Flatten)	(None, 1500)	0	clean_subcategories_Embedding[0][0]
dense_1 (Dense)	(None, 128)	384	numerical_input[0][0]
concatenate_18 (Concatenate)	(None, 28811)	0	flatten_1[0][0] flatten_2[0][0] flatten_3[0][0] flatten_4[0][0] flatten_5[0][0] flatten_6[0][0] dense_1[0][0]
dense_80 (Dense)	(None, 256)	7375872	concatenate_18[0][0]
dropout_49 (Dropout)	(None, 256)	0	dense_80[0][0]
dense_81 (Dense)	(None, 128)	32896	dropout_49[0][0]
dropout_50 (Dropout)	(None, 128)	0	dense_81[0][0]
dense_82 (Dense)	(None, 64)	8256	dropout_50[0][0]
dropout_51 (Dropout)	(None, 64)	0	dense_82[0][0]
dense_83 (Dense)	(None, 2)	130	dropout_51[0][0]
=====			
Total params: 21,358,861			
Trainable params: 7,659,061			
Non-trainable params: 13,699,800			

In [143]:

```
from keras.utils import plot_model
plot_model(model, to_file='model_1.png', show_shapes=True, show_layer_names=True)
```

Out[143]:





In [144]:

```
history = model.fit(x=[X1_tr,X2_tr,X3_tr,X4_tr, X5_tr, X6_tr,X7_tr], y=y_train, validation_data=([X1_cv,X2_cv,X3_cv,X4_cv, X5_cv, X6_cv,X7_cv],y_cv),epochs=10,batch_size=300,verbose=2)
```

Train on 64000 samples, validate on 16000 samples

```
Epoch 1/10
- 100s - loss: 0.3608 - auroc: 0.7872 - val_loss: 0.3691 - val_auroc: 0.7487
Epoch 2/10
- 93s - loss: 0.3380 - auroc: 0.7965 - val_loss: 0.3683 - val_auroc: 0.7504
Epoch 3/10
- 91s - loss: 0.3344 - auroc: 0.7991 - val_loss: 0.3689 - val_auroc: 0.7497
Epoch 4/10
- 93s - loss: 0.3311 - auroc: 0.8028 - val_loss: 0.3729 - val_auroc: 0.7505
Epoch 5/10
- 93s - loss: 0.3274 - auroc: 0.8054 - val_loss: 0.3686 - val_auroc: 0.7496
Epoch 6/10
- 92s - loss: 0.3230 - auroc: 0.8095 - val_loss: 0.3722 - val_auroc: 0.7488
Epoch 7/10
- 95s - loss: 0.3189 - auroc: 0.8114 - val_loss: 0.3690 - val_auroc: 0.7478
Epoch 8/10
- 95s - loss: 0.3145 - auroc: 0.8131 - val_loss: 0.3732 - val_auroc: 0.7468
Epoch 9/10
- 93s - loss: 0.3085 - auroc: 0.8172 - val_loss: 0.3751 - val_auroc: 0.7466
Epoch 10/10
- 93s - loss: 0.3027 - auroc: 0.8201 - val_loss: 0.3809 - val_auroc: 0.7449
```

In [0]:

```
score = model.evaluate(x=[X1_test,X2_test,X3_test,X4_test, X5_test, X6_test, X7_test], y=y_test, verbose=2,batch_size=500)
```

In [146]:

```
print("Test Loss:", score[0])
print("Test AUC:", score[1])
```

```
Test Loss: 0.38501315861940383
Test AUC: 0.7441678294901292
```

In [147]:



```
plt.plot(history.history['auroc'])
plt.plot(history.history['val_auroc'])

plt.title('model auc')
plt.ylabel('auc')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])

plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
```

```
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



In [148]:


```
# serialize weights to HDF5
model.save_weights("model_1.h5")
print("Saved model to disk")
```


Saved model to disk

Model 2

In [149]:

```
from sklearn.feature_extraction.text import TfidfVectorizer
import seaborn as sns
vectorizer = TfidfVectorizer()
vectorizer.fit(X_train['essay'].values)
plt.boxplot(list(vectorizer.idf_))
plt.xlabel('Essay')
plt.ylabel('IDF Value')
plt.show()
```





In [150]:

```
tenth_percentile=np.quantile((vectorizer.idf_),0.10)
ninty_percentile=np.quantile((vectorizer.idf_),0.90)
print(tenth_percentile)
print(ninty_percentile)
dictionary = dict(zip(vectorizer.get_feature_names(), list(vectorizer.idf_)))
filtered_words=[]
for k,v in dictionary.items():
    if v > tenth_percentile and v < ninty_percentile:
        filtered_words.append(k)

len(filtered_words)
```

7.375306104990596
11.373506806659794

Out[150]:

23474

In [0]:

```
tokenizer = Tokenizer()
tokenizer.fit_on_texts(filtered_words)

X8_tr = np.array(tokenizer.texts_to_sequences(X_train['essay'].values))
X8_cv = np.array(tokenizer.texts_to_sequences(X_cv['essay'].values))
X8_test = np.array(tokenizer.texts_to_sequences(X_test['essay'].values))
```

In [0]:

```
vocab_size = len(tokenizer.word_index) + 1

maxlen = 200

X8_tr = pad_sequences(X8_tr, padding='post', maxlen=maxlen)
X8_cv = pad_sequences(X8_cv, padding='post', maxlen=maxlen)
X8_test = pad_sequences(X8_test, padding='post', maxlen=maxlen)
```

In [42]:

```
print(X8_tr.shape)
print(X8_cv.shape)
print(X8_test.shape)
```

(64000, 200)
(16000, 200)
(20000, 200)

In [0]:

```
with open('glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())
```

In [0]:

```
embeddings_dictionary = dict()
for word in glove_words:
    vector_dimensions = model[word]
    embeddings_dictionary[word] = vector_dimensions
```

In [0]:

```
embedding_matrix = np.zeros((vocab_size, 300))
for word, index in tokenizer.word_index.items():
    embedding_vector = embeddings_dictionary.get(word)
    if embedding_vector is not None:
        embedding_matrix[index] = embedding_vector
```

In [46]:

```
embedding_matrix.shape
```

Out[46]:

```
(23475, 300)
```

In [156]:

```
input_8 = Input(shape=(maxlen,), name='essay_tfidf_input')
print(input_8.shape)
input_8_embedding = Embedding(vocab_size, 300, weights=[embedding_matrix], trainable=False)(input_8)
print(input_8_embedding.shape)
input_8_lstm = LSTM(128, return_sequences=True)(input_8_embedding)
print(input_8_lstm.shape)
input_8_flatten = Flatten()(input_8_lstm)
print(input_8_flatten.shape)
```

```
(?, 200)
(?, 200, 300)
(?, ?, 128)
(?, ?)
```

In [0]:

```
#At the end we concatenate altogether and add other Dense layers
output_2 = Concatenate()
([input_8_flatten, input_2_flatten, input_3_flatten, input_4_flatten, input_5_flatten, input_6_flatten,
input_7_dense])
output_2 = Dense(256, kernel_initializer="he_uniform", activation='relu', kernel_regularizer=l2(0.001))
(output_2)
output_2 = Dropout(0.3)(output_2)
output_2 = Dense(128, kernel_initializer="he_uniform", activation='relu', kernel_regularizer=l2(0.001))
(output_2)
output_2 = Dropout(0.5)(output_2)
output_2 = Dense(64, kernel_initializer="he_uniform", activation='relu', kernel_regularizer=l2(0.001))
(output_2)
output_2 = Dense(2, activation='softmax')(output_2)
```

In [158]:

```
model_2 = Model(inputs=[input_8, input_2, input_3, input_4, input_5, input_6, input_7], outputs=output_2)
model_2.compile(loss='binary_crossentropy', optimizer=optimizers.SGD(lr=0.01, decay=1e-3, momentum=0.9), metrics=[auroc])
model_2.summary()
```

Model: "model_22"

Layer (type)	Output Shape	Param #	Connected to
essay_tfidf_input (InputLayer)	(None, 200)	0	
embedding_3 (Embedding)	(None, 200, 300)	7042500	essay_tfidf_input[0][0]
school_state_input (InputLayer)	(None, 51)	0	
teacher_prefix_input (InputLayer)	(None, 5)	0	
project_grade_category_input (InputLayer)	(None, 4)	0	
clean_categories_input (InputLayer)	(None, 9)	0	

clean_subcategories_input (InputLayer)	(None, 30)	0	
lstm_3 (LSTM)	(None, 200, 128)	219648	embedding_3[0][0]
school_state_Embedding (Embedding)	(None, 51, 26)	1326	school_state_input[0][0]
teacher_prefix_Embedding (Embedding)	(None, 5, 3)	15	teacher_prefix_input[0][0]
project_grade_category_Embedding (Embedding)	(None, 4, 2)	8	project_grade_category_input[0][0]
clean_categories_Embedding (Embedding)	(None, 9, 26)	1326	clean_categories_input[0][0]
clean_subcategories_Embedding (Embedding)	(None, 30, 50)	19200	clean_subcategories_input[0][0]
numerical_input (InputLayer)	(None, 2)	0	
flatten_9 (Flatten)	(None, 25600)	0	lstm_3[0][0]
flatten_2 (Flatten)	(None, 1326)	0	school_state_Embedding[0][0]
flatten_3 (Flatten)	(None, 15)	0	teacher_prefix_Embedding[0][0]
flatten_4 (Flatten)	(None, 8)	0	project_grade_category_Embedding[0][0]
flatten_5 (Flatten)	(None, 234)	0	clean_categories_Embedding[0][0]
flatten_6 (Flatten)	(None, 1500)	0	clean_subcategories_Embedding[0][0]
dense_1 (Dense)	(None, 128)	384	numerical_input[0][0]
concatenate_19 (Concatenate)	(None, 28811)	0	flatten_9[0][0] flatten_2[0][0] flatten_3[0][0] flatten_4[0][0] flatten_5[0][0] flatten_6[0][0] dense_1[0][0]
dense_84 (Dense)	(None, 256)	7375872	concatenate_19[0][0]
dropout_52 (Dropout)	(None, 256)	0	dense_84[0][0]
dense_85 (Dense)	(None, 128)	32896	dropout_52[0][0]
dropout_53 (Dropout)	(None, 128)	0	dense_85[0][0]
dense_86 (Dense)	(None, 64)	8256	dropout_53[0][0]
dense_87 (Dense)	(None, 2)	130	dense_86[0][0]
=====			
Total params: 14,701,561			
Trainable params: 7,659,061			
Non-trainable params: 7,042,500			

In [159]:

```
plot_model(model_2, to_file='model_2.png', show_shapes=True, show_layer_names=True)
```

Out[159]:



In [160]:

```
history = model_2.fit(x=[X8_tr,X2_tr,X3_tr,X4_tr, X5_tr, X6_tr,X7_tr], y=y_train, validation_data=(
[X8_cv,X2_cv,X3_cv,X4_cv, X5_cv, X6_cv,X7_cv],y_cv),epochs=10,batch_size=500,verbose=2)
```

```
Train on 64000 samples, validate on 16000 samples
Epoch 1/10
- 68s - loss: 1.3157 - auroc: 0.5545 - val_loss: 1.2714 - val_auroc: 0.6000
Epoch 2/10
- 62s - loss: 1.2558 - auroc: 0.5857 - val_loss: 1.2340 - val_auroc: 0.6114
Epoch 3/10
- 62s - loss: 1.2197 - auroc: 0.6003 - val_loss: 1.2017 - val_auroc: 0.6154
Epoch 4/10
- 62s - loss: 1.1893 - auroc: 0.6084 - val_loss: 1.1735 - val_auroc: 0.6174
Epoch 5/10
- 63s - loss: 1.1628 - auroc: 0.6127 - val_loss: 1.1490 - val_auroc: 0.6195
Epoch 6/10
- 62s - loss: 1.1384 - auroc: 0.6184 - val_loss: 1.1267 - val_auroc: 0.6210
Epoch 7/10
- 63s - loss: 1.1171 - auroc: 0.6215 - val_loss: 1.1063 - val_auroc: 0.6218
Epoch 8/10
- 63s - loss: 1.0979 - auroc: 0.6232 - val_loss: 1.0880 - val_auroc: 0.6230
Epoch 9/10
- 62s - loss: 1.0797 - auroc: 0.6278 - val_loss: 1.0717 - val_auroc: 0.6234
Epoch 10/10
- 63s - loss: 1.0633 - auroc: 0.6322 - val_loss: 1.0566 - val_auroc: 0.6228
```

In [0]:

```
score = model_2.evaluate(x=[X8_test,X2_test,X3_test,X4_test, X5_test, X6_test, X7_test], y=y_test,
verbose=2,batch_size=500)
```

In [162]:

```
print("Test Loss:", score[0])
print("Test AUC:", score[1])
```

```
Test Loss: 1.0564408168196677
Test AUC: 0.6289043083580907
```


In [163]:



```
plt.plot(history.history['auroc'])
plt.plot(history.history['val_auroc'])

plt.title('model auc')
plt.ylabel('auc')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])

plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```





In [164]:

```
# serialize weights to HDF5
model_2.save_weights("model_2.h5")
print("Saved model to disk")
```

Saved model to disk

Model 3

In [165]:

```
from sklearn.preprocessing import StandardScaler
scalar = StandardScaler()
X9_tr = scalar.fit_transform(X_train['price'].values.reshape(-1,1)) # finding the mean and standard
deviation of this data
X9_cv = scalar.transform(X_cv['price'].values.reshape(-1,1))
X9_test = scalar.transform(X_test['price'].values.reshape(-1,1))
print(X9_tr.shape)
print(X9_cv.shape)
print(X9_test.shape)
```

```
(64000, 1)
(16000, 1)
(20000, 1)
```

In [166]:

```
scalar = StandardScaler()
X10_tr = scalar.fit_transform(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1, 1))
X10_cv = scalar.transform(X_cv['teacher_number_of_previously_posted_projects'].values.reshape(-1, 1))
X10_test = scalar.transform(X_test['teacher_number_of_previously_posted_projects'].values.reshape(-1, 1))
print(X10_tr.shape)
print(X10_cv.shape)
print(X10_test.shape)
```

```
(64000, 1)
(16000, 1)
(20000, 1)
```

In [167]:

In [167]:

```
numeric_tr = np.hstack((X2_tr,X3_tr,X4_tr,X5_tr,X6_tr,X9_tr,X10_tr))
numeric_cv = np.hstack((X2_cv,X3_cv,X4_cv,X5_cv,X6_cv,X9_cv,X10_cv))
numeric_test = np.hstack((X2_test,X3_test,X4_test,X5_test,X6_test,X9_test,X10_test))
print(numeric_tr.shape)
print(numeric_cv.shape)
print(numeric_test.shape)
```

```
(64000, 101)
(16000, 101)
(20000, 101)
```

In [168]:

```
numeric_tr=np.expand_dims(numeric_tr,axis=2)
numeric_cv=np.expand_dims(numeric_cv,axis=2)
numeric_test=np.expand_dims(numeric_test,axis=2)
print(numeric_tr.shape)
print(numeric_cv.shape)
print(numeric_test.shape)
```

```
(64000, 101, 1)
(16000, 101, 1)
(20000, 101, 1)
```

In [169]:

```
input_9 = Input(shape=(numeric_tr.shape[1],numeric_tr.shape[2]),name='combined_input')
print(input_9.shape)
```

```
(?, 101, 1)
```

In [0]:

#At the end we concatenate altogether and add other Dense layers

```
#output_3=tf.reshape(output_3,[-1,output_3.shape[1],output_3.shape[1]])
#print(output_3.shape)
```

```
output_3 = Conv1D(128, 5, strides=1,activation="relu",kernel_initializer='he_uniform',padding='same') (input_9)
output_3 = MaxPooling1D(pool_size=5) (output_3)
output_3= Dropout(0.2) (output_3)
output_3 = Conv1D(64, 5, activation="relu") (output_3)
output_3 = MaxPooling1D(pool_size=5) (output_3)
output_3= Dropout(0.4) (output_3)
output_3 = Flatten() (output_3)
```

```
output_4 = Concatenate() ([input_1_flatten,output_3])
output_4 = Dense(256, kernel_initializer="he_uniform",activation='relu') (output_4)
output_4= Dropout(0.2) (output_4)
output_4 = Dense(128, kernel_initializer="he_uniform",activation='relu') (output_4)
output_4= Dropout(0.3) (output_4)
output_4 = Dense(64, kernel_initializer="he_uniform", activation='relu') (output_4)
output_4= Dropout(0.4) (output_4)
output_4 = Dense(2, activation='softmax') (output_4)
```

In [171]:

```
from keras import optimizers
model_3 = Model(inputs=[input_1,input_2,input_3,input_4,input_5,input_6,input_9], outputs=output_4)
model_3.compile(loss='binary_crossentropy', optimizer=optimizers.SGD(lr=0.01, decay=1e-6, momentum=0.9),metrics=[auc])
model_3.summary()
```

Model: "model_23"

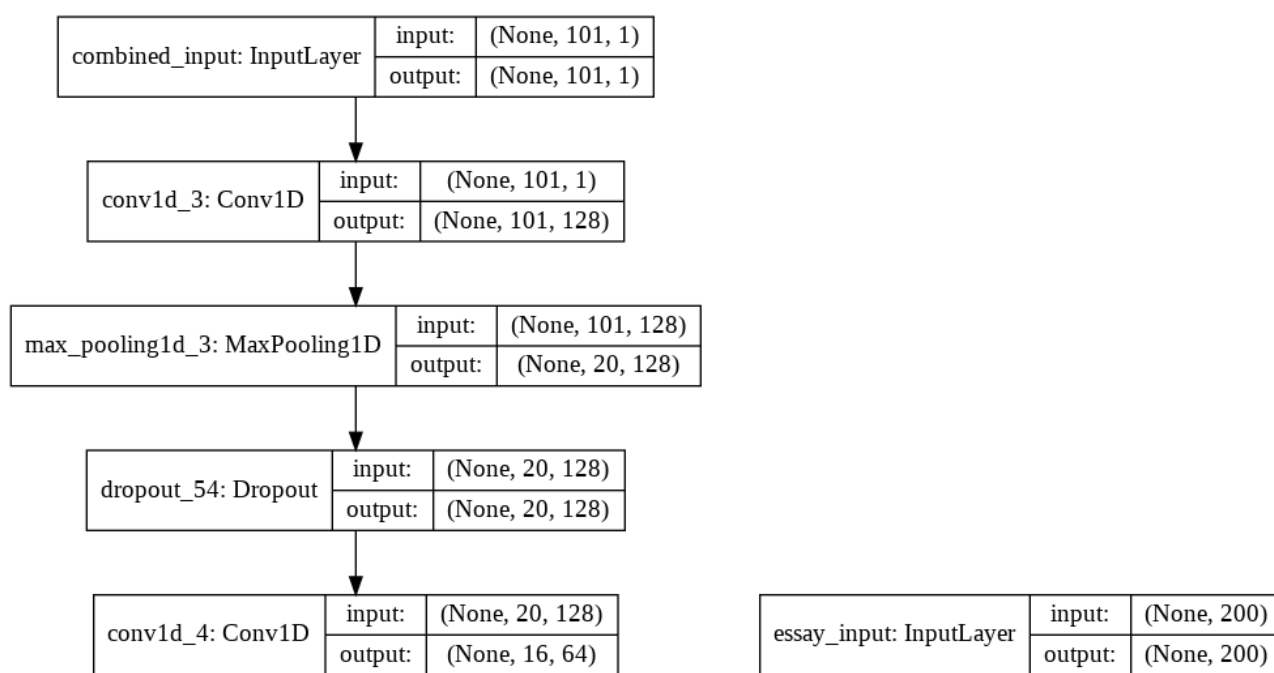
Layer (type)	Output Shape	Param #	Connected to
=====			
combined_input (InputLayer)	(None, 101, 1)	0	

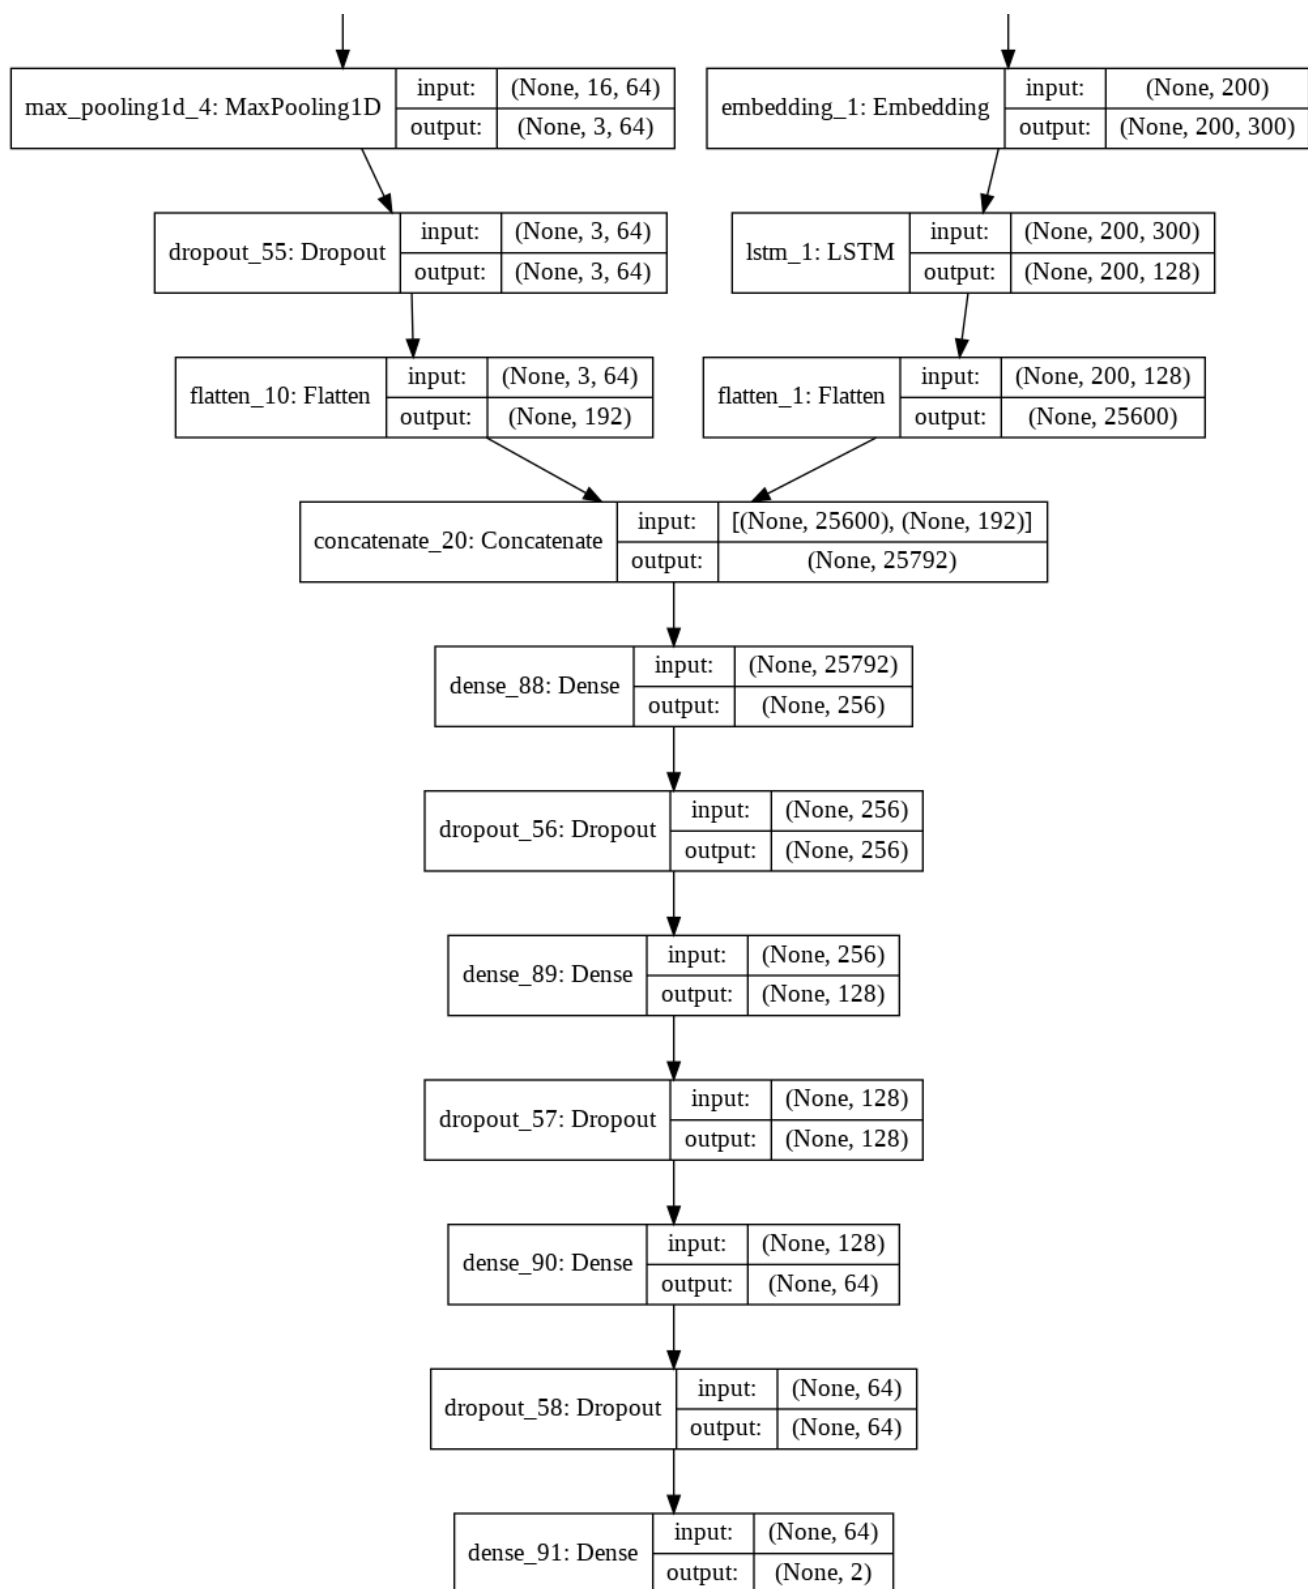
conv1d_3 (Conv1D)	(None, 101, 128)	768	combined_input[0][0]
max_pooling1d_3 (MaxPooling1D)	(None, 20, 128)	0	conv1d_3[0][0]
dropout_54 (Dropout)	(None, 20, 128)	0	max_pooling1d_3[0][0]
essay_input (InputLayer)	(None, 200)	0	
conv1d_4 (Conv1D)	(None, 16, 64)	41024	dropout_54[0][0]
embedding_1 (Embedding)	(None, 200, 300)	13699800	essay_input[0][0]
max_pooling1d_4 (MaxPooling1D)	(None, 3, 64)	0	conv1d_4[0][0]
lstm_1 (LSTM)	(None, 200, 128)	219648	embedding_1[0][0]
dropout_55 (Dropout)	(None, 3, 64)	0	max_pooling1d_4[0][0]
flatten_1 (Flatten)	(None, 25600)	0	lstm_1[0][0]
flatten_10 (Flatten)	(None, 192)	0	dropout_55[0][0]
concatenate_20 (Concatenate)	(None, 25792)	0	flatten_1[0][0] flatten_10[0][0]
dense_88 (Dense)	(None, 256)	6603008	concatenate_20[0][0]
dropout_56 (Dropout)	(None, 256)	0	dense_88[0][0]
dense_89 (Dense)	(None, 128)	32896	dropout_56[0][0]
dropout_57 (Dropout)	(None, 128)	0	dense_89[0][0]
dense_90 (Dense)	(None, 64)	8256	dropout_57[0][0]
dropout_58 (Dropout)	(None, 64)	0	dense_90[0][0]
dense_91 (Dense)	(None, 2)	130	dropout_58[0][0]
=====			
Total params: 20,605,530			
Trainable params: 6,905,730			
Non-trainable params: 13,699,800			

In [172]:

```
plot_model(model_3, to_file='model_3.png', show_shapes=True, show_layer_names=True)
```

Out[172]:





In [173]:

```
history = model_3.fit(x=[X8_tr,X2_tr,X3_tr,X4_tr, X5_tr, X6_tr,numeric_tr], y=y_train, validation_data=([X8_cv,X2_cv,X3_cv,X4_cv, X5_cv, X6_cv,numeric_cv],y_cv),epochs=10,batch_size=500,verbose=2)
```

Train on 64000 samples, validate on 16000 samples

Epoch 1/10

- 71s - loss: 0.4506 - auroc: 0.5017 - val_loss: 0.4343 - val_auroc: 0.5238

Epoch 2/10

- 64s - loss: 0.4305 - auroc: 0.5098 - val_loss: 0.4311 - val_auroc: 0.5359

Epoch 3/10

- 62s - loss: 0.4275 - auroc: 0.5159 - val_loss: 0.4264 - val_auroc: 0.5296

Epoch 4/10

- 63s - loss: 0.4255 - auroc: 0.5255 - val_loss: 0.4246 - val_auroc: 0.5280

Epoch 5/10

- 63s - loss: 0.4249 - auroc: 0.5268 - val_loss: 0.4248 - val_auroc: 0.5323

```
Epoch 6/10
- 62s - loss: 0.4242 - auroc: 0.5316 - val_loss: 0.4251 - val_auroc: 0.5379
Epoch 7/10
- 62s - loss: 0.4232 - auroc: 0.5445 - val_loss: 0.4238 - val_auroc: 0.5442
Epoch 8/10
- 62s - loss: 0.4230 - auroc: 0.5458 - val_loss: 0.4236 - val_auroc: 0.5453
Epoch 9/10
- 62s - loss: 0.4229 - auroc: 0.5496 - val_loss: 0.4233 - val_auroc: 0.5477
Epoch 10/10
- 63s - loss: 0.4227 - auroc: 0.5522 - val_loss: 0.4235 - val_auroc: 0.5384
```

In [0]:

```
score = model_3.evaluate(x=[X8_test,X2_test,X3_test,X4_test, X5_test, X6_test, numeric_test], y=y_test, verbose=2,batch_size=500)
```

In [175]:

```
print("Test Loss:", score[0])
print("Test AUC:", score[1])
```

```
Test Loss: 0.4226397812366486
Test AUC: 0.5524177426764195
```



In [176]:


```
plt.plot(history.history['auroc'])
plt.plot(history.history['val_auroc'])

plt.title('model auc')
plt.ylabel('auc')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])

plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```





In [177]:

```
# serialize weights to HDF5
model_3.save_weights("model_3.h5")
print("Saved model to disk")
```

Saved model to disk