

DonorsChoose

DonorsChoose.org receives hundreds of thousands of project proposals each year for classroom projects in need of funding. Right now, a large number of volunteers is needed to manually screen each submission before it's approved to be posted on the DonorsChoose.org website.

Next year, DonorsChoose.org expects to receive close to 500,000 project proposals. As a result, there are three main problems they need to solve:

- How to scale current manual processes and resources to screen 500,000 projects so that they can be posted as quickly and as efficiently as possible
- How to increase the consistency of project vetting across different volunteers to improve the experience for teachers
- How to focus volunteer time on the applications that need the most assistance

The goal of the competition is to predict whether or not a DonorsChoose.org project proposal submitted by a teacher will be approved, using the text of project descriptions as well as additional metadata about the project, teacher, and school. DonorsChoose.org can then use this information to identify projects most likely to need further review before approval.

About the DonorsChoose Data Set

The `train.csv` data set provided by DonorsChoose contains the following features:

Feature	Description
<code>project_id</code>	A unique identifier for the proposed project. Example: p036502
<code>project_title</code>	Title of the project. Examples: Art Will Make You Happy! First Grade Fun
<code>project_grade_category</code>	Grade level of students for which the project is targeted. One of the following enumerated values: Grades PreK-2 Grades 3-5 Grades 6-8 Grades 9-12
<code>project_subject_categories</code>	One or more (comma-separated) subject categories for the project from the following enumerated list of values: Applied Learning Care & Hunger Health & Sports History & Civics Literacy & Language Math & Science Music & The Arts Special Needs Warmth Examples: Music & The Arts Literacy & Language, Math & Science
<code>school_state</code>	State where school is located (Two-letter U.S. postal code). Example: WY
<code>project_subject_subcategories</code>	One or more (comma-separated) subject subcategories for the project. Examples: Literacy Literature & Writing, Social Sciences
<code>project_resource_summary</code>	An explanation of the resources needed for the project. Example: My students need hands on literacy materials to manage sensory needs!
<code>project_essay_1</code>	First application essay*
<code>project_essay_2</code>	Second application essay*
<code>project_essay_3</code>	Third application essay*

Feature	Description
project_essay_4	Fourth application essay
project_submitted_datetime	Datetime when project application was submitted. Example: 2016-04-28 12:43:56.245
teacher_id	A unique identifier for the teacher of the proposed project. Example: bdf8baa8fedef6bfeec7ae4ff1c15c56
teacher_prefix	Teacher's title. One of the following enumerated values: <ul style="list-style-type: none"> nan Dr. Mr. Mrs. Ms. Teacher.
teacher_number_of_previously_posted_projects	Number of project applications previously submitted by the same teacher. Example: 2

* See the section **Notes on the Essay Data** for more details about these features.

Additionally, the `resources.csv` data set provides more data about the resources required for each project. Each line in this file represents a resource required by a project:

Feature	Description
id	A <code>project_id</code> value from the <code>train.csv</code> file. Example: p036502
description	Description of the resource. Example: Tenor Saxophone Reeds, Box of 25
quantity	Quantity of the resource required. Example: 3
price	Price of the resource required. Example: 9.95

Note: Many projects require multiple resources. The `id` value corresponds to a `project_id` in `train.csv`, so you use it as a key to retrieve all resources needed for a project:

The data set contains the following label (the value you will attempt to predict):

Label	Description
project_is_approved	A binary flag indicating whether DonorsChoose approved the project. A value of 0 indicates the project was not approved, and a value of 1 indicates the project was approved.

Notes on the Essay Data

Prior to May 17, 2016, the prompts for the essays were as follows:

- __project_essay_1__: "Introduce us to your classroom"
- __project_essay_2__: "Tell us more about your students"
- __project_essay_3__: "Describe how your students will use the materials you're requesting"
- __project_essay_3__: "Close by sharing why your project will make a difference"

Starting on May 17, 2016, the number of essays was reduced from 4 to 2, and the prompts for the first 2 essays were changed to the following:

- __project_essay_1__: "Describe your students: What makes your students special? Specific details about their background, your neighborhood, and your school are all helpful."
- __project_essay_2__: "About your project: How will these materials make a difference in your students' learning and improve their school lives?"

For all projects with `project_submitted_datetime` of 2016-05-17 and later, the values of `project_essay_3` and `project_essay_4` will be NaN.

In [0]:

```
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
warnings.simplefilter("ignore")
warnings.warn("ignore")
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
```

```

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn import model_selection
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

```

In [0]:

```

from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve, auc
from sklearn import preprocessing

```

In [3]:

```

from keras.preprocessing.text import one_hot
from keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers.core import Activation, Dropout, Dense, Reshape
from keras.layers import Flatten, LSTM, Lambda
from keras.models import Model
from keras.layers.embeddings import Embedding
from keras.preprocessing.text import Tokenizer
from keras.layers import Input
from keras.layers import Concatenate
from keras.utils import to_categorical
from keras.layers import Conv1D, MaxPooling1D
from keras import regularizers
from keras.layers.normalization import BatchNormalization
from keras.utils import plot_model
from keras import optimizers

```

Using TensorFlow backend.

The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x.

We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the `%tensorflow_version 1.x` magic: [more info](#).

1.1 Reading Data

In [0]:

```

#https://stackabuse.com/python-for-nlp-creating-multi-data-type-classification-models-with-keras/
#https://www.pyimagesearch.com/2019/01/21/regression-with-keras/
#https://github.com/mmortazavi/EntityEmbedding-Working_Example/blob/master/EntityEmbedding.ipynb
#https://www.pyimagesearch.com/2019/02/04/keras-multiple-inputs-and-mixed-data/
#https://machinelearningmastery.com/cnn-models-for-human-activity-recognition-time-series-classification/

```

In [0]:

In [0]:

```
preprocessed_data = pd.read_csv('preprocessed_data.csv')
```

In [5]:

```
print("Number of data points in preprocessed data", preprocessed_data.shape)
```

Number of data points in preprocessed data (109248, 9)

In [6]:

```
preprocessed_data=preprocessed_data.sample(n=100000)
preprocessed_data.head()
```

Out[6]:

	school_state	teacher_prefix	project_grade_category	teacher_number_of_previously_posted_projects	project_is_approved	clean_
69638	in	mrs	grades_prek_2	1	1	he
7741	wa	mrs	grades_3_5	5	1	appl literacy
45671	ak	mrs	grades_3_5	2	0	appl
52272	wi	mrs	grades_prek_2	48	1	appl
23037	wa	mrs	grades_prek_2	2	1	literacy

In [0]:

```
X=preprocessed_data.drop(columns=['project_is_approved'],axis=1)
y=preprocessed_data['project_is_approved']
```

In [0]:

```
label_encoder = preprocessing.LabelEncoder()
y = label_encoder.fit_transform(y)
```

In [0]:

```
X_1, X_test, y_1, y_test = model_selection.train_test_split(X, y, test_size=0.2, random_state=0, stratify=y)

# split the train data set into cross validation train and cross validation test
X_train, X_cv, y_train, y_cv = model_selection.train_test_split(X_1, y_1, test_size=0.2, random_state=0, stratify=y_1)
```

In [0]:

```
y_train = to_categorical(y_train)
y_cv = to_categorical(y_cv)
y_test = to_categorical(y_test)
```

In [0]:

```
tokenizer = Tokenizer()
tokenizer.fit_on_texts(X_train['essay'].values)

X1_tr = np.array(tokenizer.texts_to_sequences(X_train['essay'].values))
X1_cv = np.array(tokenizer.texts_to_sequences(X_cv['essay'].values))
X1_test = np.array(tokenizer.texts_to_sequences(X_test['essay'].values))
```

In [0]:

```
vocab_size = len(tokenizer.word_index) + 1

maxlen = 200

X1_tr = pad_sequences(X1_tr, padding='post', maxlen=maxlen)
X1_cv = pad_sequences(X1_cv, padding='post', maxlen=maxlen)
X1_test = pad_sequences(X1_test, padding='post', maxlen=maxlen)
```

In [13]:

```
print(X1_tr.shape)
print(X1_cv.shape)
print(X1_test.shape)
```

```
(64000, 200)
(16000, 200)
(20000, 200)
```

In [0]:

```
with open('glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())
```

In [0]:

```
embeddings_dictionary = dict()
for word in glove_words:
    vector_dimensions = model[word]
    embeddings_dictionary[word] = vector_dimensions
```

In [0]:

```
embedding_matrix = np.zeros((vocab_size, 300))
for word, index in tokenizer.word_index.items():
    embedding_vector = embeddings_dictionary.get(word)
    if embedding_vector is not None:
        embedding_matrix[index] = embedding_vector
```

In [18]:

```
embedding_matrix.shape
```

Out[18]:

```
(45838, 300)
```

In [19]:

```
input_1 = Input(shape=(maxlen,), name='essay_input')
print(input_1.shape)
input_1_embedding = Embedding(vocab_size, 300, weights=[embedding_matrix], trainable=False)(input_1)
print(input_1_embedding.shape)
input_1_lstm = LSTM(128, return_sequences=True)(input_1_embedding)
print(input_1_lstm.shape)
input_1_flatten = Flatten()(input_1_lstm)
print(input_1_flatten.shape)
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.
```

```
(?, 200)
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:190: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:203: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:207: The name tf.global_variables is deprecated. Please use tf.compat.v1.global_variables instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:216: The name tf.is_variable_initialized is deprecated. Please use tf.compat.v1.is_variable_initialized instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:223: The name tf.variables_initializer is deprecated. Please use tf.compat.v1.variables_initializer instead.
```

```
(?, 200, 300)
```

```
(?, ?, 128)
```

```
(?, ?)
```

```
In [0]:
```

```
categoricals=['school_state','teacher_prefix','project_grade_category','clean_categories','clean_subcategories']
numericals=['teacher_number_of_previously_posted_projects','price']
```

```
In [0]:
```

```
embed_cols=[i for i in X_train[categoricals]]

for i in embed_cols:
    print(i,X_train[i].nunique())
```

```
In [22]:
```

```
tokenizer = Tokenizer()
tokenizer.fit_on_texts(X_train['school_state'].values)

X2_tr = np.array(tokenizer.texts_to_sequences(X_train['school_state'].values))
X2_cv = np.array(tokenizer.texts_to_sequences(X_cv['school_state'].values))
X2_test = np.array(tokenizer.texts_to_sequences(X_test['school_state'].values))
cat_emb_name= 'school_state_Embedding'
no_of_unique_cat = X_train['school_state'].nunique()
embedding_size = int(min(np.ceil((no_of_unique_cat)/2), 50 ))
input_2 = Input(shape=(1,),name='school_state_input')
print(input_2.shape)
input_2_embedding = Embedding(no_of_unique_cat+1, embedding_size,input_length=1, name=cat_emb_name)(input_2)
print(input_2_embedding.shape)
input_2_flatten=Flatten()(input_2_embedding)
print(input_2_flatten.shape)
```

```
(?, 1)
(?, 1, 26)
(?, ?)
```

In [23]:

```
tokenizer = Tokenizer()
tokenizer.fit_on_texts(X_train['teacher_prefix'].values)

X3_tr = np.array(tokenizer.texts_to_sequences(X_train['teacher_prefix'].values))
X3_cv = np.array(tokenizer.texts_to_sequences(X_cv['teacher_prefix'].values))
X3_test = np.array(tokenizer.texts_to_sequences(X_test['teacher_prefix'].values))
cat_emb_name= 'teacher_prefix_Embedding'
no_of_unique_cat = X_train['teacher_prefix'].nunique()
embedding_size = int(min(np.ceil((no_of_unique_cat)/2), 50 ))
input_3 = Input(shape=(1,),name='teacher_prefix_input')
print(input_3.shape)
input_3_embedding = Embedding(no_of_unique_cat+1, embedding_size,input_length=1, name=cat_emb_name)
(input_3)
print(input_3_embedding.shape)
input_3_flatten=Flatten()(input_3_embedding)
print(input_3_flatten.shape)

(?, 1)
(?, 1, 3)
(?, ?)
```

In [0]:

```
#X_train['project_grade_category']=X_train['project_grade_category'].str.replace('_', '')
#X_cv['project_grade_category']=X_cv['project_grade_category'].str.replace('_', '')
#X_test['project_grade_category']=X_test['project_grade_category'].str.replace('_', '')
```

In [24]:

```
tokenizer = Tokenizer(filters='!"#$%&()*+,-./:;<=>?@[\\]^`{|}~\t\n')
tokenizer.fit_on_texts(X_train['project_grade_category'].values)

X4_tr = np.array(tokenizer.texts_to_sequences(X_train['project_grade_category'].values))
X4_cv = np.array(tokenizer.texts_to_sequences(X_cv['project_grade_category'].values))
X4_test = np.array(tokenizer.texts_to_sequences(X_test['project_grade_category'].values))
cat_emb_name= 'project_grade_category_Embedding'
no_of_unique_cat = X_train['project_grade_category'].nunique()
embedding_size = int(min(np.ceil((no_of_unique_cat)/2), 50 ))
input_4 = Input(shape=(1,),name='project_grade_category_input')
print(input_4.shape)
input_4_embedding = Embedding(no_of_unique_cat+1, embedding_size,input_length=1,name=cat_emb_name)(
input_4)
print(input_4_embedding.shape)
input_4_flatten=Flatten()(input_4_embedding)
print(input_4_flatten.shape)

(?, 1)
(?, 1, 2)
(?, ?)
```

In [25]:

```
tokenizer = Tokenizer(filters='!"#$%&()*+,-./:;<=>?@[\\]^`{|}~\t\n')
tokenizer.fit_on_texts(X_train['clean_categories'].values)

X5_tr = np.array(tokenizer.texts_to_sequences(X_train['clean_categories'].values))
X5_cv = np.array(tokenizer.texts_to_sequences(X_cv['clean_categories'].values))
X5_test = np.array(tokenizer.texts_to_sequences(X_test['clean_categories'].values))
maxlen=max(len(l) for l in X5_tr)
X5_tr = pad_sequences(X5_tr, padding='post', maxlen=maxlen)
X5_cv = pad_sequences(X5_cv, padding='post', maxlen=maxlen)
X5_test = pad_sequences(X5_test, padding='post', maxlen=maxlen)
cat_emb_name= 'clean_categories_Embedding'
no_of_unique_cat = X_train['clean_categories'].nunique()
embedding_size = int(min(np.ceil((no_of_unique_cat)/2), 50 ))
input_5 = Input(shape=(maxlen,),name='clean categories input')
```

```

print(input_5.shape)
input_5_embedding = Embedding(no_of_unique_cat+1,
embedding_size,input_length=maxlen,name=cat_emb_name)(input_5)
print(input_5_embedding.shape)
input_5_flatten=Flatten()(input_5_embedding)
print(input_5_flatten.shape)

```

```

(?, 3)
(?, 3, 26)
(?, ?)

```

In [26]:

```

tokenizer = Tokenizer(filters='!"#$%&()*+,-/:;<=>?@[\\]^_`{|}~\t\n')
tokenizer.fit_on_texts(X_train['clean_subcategories'].values)

X6_tr = np.array(tokenizer.texts_to_sequences(X_train['clean_subcategories'].values))
X6_cv = np.array(tokenizer.texts_to_sequences(X_cv['clean_subcategories'].values))
X6_test = np.array(tokenizer.texts_to_sequences(X_test['clean_subcategories'].values))
maxlen=max(len(l) for l in X6_tr)
X6_tr = pad_sequences(X6_tr, padding='post', maxlen=maxlen)
X6_cv = pad_sequences(X6_cv, padding='post', maxlen=maxlen)
X6_test = pad_sequences(X6_test, padding='post', maxlen=maxlen)
cat_emb_name= 'clean_subcategories_Embedding'
no_of_unique_cat = X_train['clean_subcategories'].nunique()
embedding_size = int(min(np.ceil((no_of_unique_cat)/2), 50 ))
input_6 = Input(shape=(maxlen,),name='clean_subcategories_input')
print(input_6.shape)
input_6_embedding = Embedding(no_of_unique_cat+1, embedding_size,input_length=maxlen, name=cat_emb_
name)(input_6)
print(input_6_embedding.shape)
input_6_flatten=Flatten()(input_6_embedding)
print(input_6_flatten.shape)

```

```

(?, 3)
(?, 3, 50)
(?, ?)

```

In [27]:

```

X7_tr = preprocessing.normalize(X_train[['teacher_number_of_previously_posted_projects', 'price']]
)
X7_cv = preprocessing.normalize(X_cv[['teacher_number_of_previously_posted_projects', 'price']]
X7_test = preprocessing.normalize(X_test[['teacher_number_of_previously_posted_projects', 'price']
])
input_7 = Input(shape=(len(X_train[numericals].columns),),name='numerical_input')
print(input_7.shape)
input_7_dense = Dense(128)(input_7)
print(input_7_dense.shape)

```

```

(?, 2)
(?, 128)

```

In [28]:

```

print(X7_tr.shape)
print(X7_cv.shape)
print(X7_test.shape)

```

```

(64000, 2)
(16000, 2)
(20000, 2)

```

In [0]:

```

#At the end we concatenate altogether and add other Dense layers
output_1 = Concatenate()
([input_1_flatten,input_2_flatten,input_3_flatten,input_4_flatten,input_5_flatten,input_6_flatten,
input_7_dense])
output_1 = Dense(512,activation='relu',kernel_initializer='he_uniform')(output_1)

```



```

output_1= Dropout(0.2)(output_1)
output_1 = Dense(256,activation='relu',kernel_initializer='he_uniform')(output_1)
output_1= Dropout(0.3)(output_1)
output_1 = Dense(128,activation='relu',kernel_initializer='he_uniform')(output_1)
output_1= Dropout(0.4)(output_1)
output_1 = Dense(2, activation='softmax')(output_1)

```

In [0]:

```

#https://stackoverflow.com/questions/41032551/how-to-compute-receiving-operating-characteristic-ro
c-and-auc-in-keras
import tensorflow as tf
def auroc(y_true, y_pred):
    return tf.py_func(roc_auc_score, (y_true, y_pred), tf.double)

```

In [58]:

```

model = Model(inputs=[input_1,input_2,input_3,input_4,input_5,input_6,input_7], outputs=output_1)
model.compile(loss='binary_crossentropy', optimizer=optimizers.SGD(lr=0.001) ,metrics=[auroc])
model.summary()

```

Model: "model_10"

Layer (type)	Output Shape	Param #	Connected to
essay_input (InputLayer)	(None, 200)	0	
embedding_1 (Embedding)	(None, 200, 300)	13751400	essay_input[0][0]
school_state_input (InputLayer)	(None, 1)	0	
teacher_prefix_input (InputLayer)	(None, 1)	0	
project_grade_category_input (InputLayer)	(None, 1)	0	
clean_categories_input (InputLayer)	(None, 3)	0	
clean_subcategories_input (InputLayer)	(None, 3)	0	
lstm_1 (LSTM)	(None, 200, 128)	219648	embedding_1[0][0]
school_state_Embedding (Embedding)	(None, 1, 26)	1352	school_state_input[0][0]
teacher_prefix_Embedding (Embedding)	(None, 1, 3)	18	teacher_prefix_input[0][0]
project_grade_category_Embedding (Embedding)	(None, 1, 2)	10	project_grade_category_input[0][0]
clean_categories_Embedding (Embedding)	(None, 3, 26)	1352	clean_categories_input[0][0]
clean_subcategories_Embedding (Embedding)	(None, 3, 50)	19650	clean_subcategories_input[0][0]
numerical_input (InputLayer)	(None, 2)	0	
flatten_1 (Flatten)	(None, 25600)	0	lstm_1[0][0]
flatten_2 (Flatten)	(None, 26)	0	school_state_Embedding[0][0]
flatten_3 (Flatten)	(None, 3)	0	teacher_prefix_Embedding[0][0]
flatten_4 (Flatten)	(None, 2)	0	project_grade_category_Embedding[0][0]
flatten_5 (Flatten)	(None, 78)	0	clean_categories_Embedding[0][0]
flatten_6 (Flatten)	(None, 150)	0	clean_subcategories_Embedding[0][0]
dense_1 (Dense)	(None, 128)	384	numerical_input[0][0]
concatenate_7 (Concatenate)	(None, 25987)	0	flatten_1[0][0] flatten_2[0][0] flatten_3[0][0] flatten_4[0][0] flatten_5[0][0] flatten_6[0][0] dense_1[0][0]

dense_26 (Dense)	(None, 512)	13305856	concatenate_7[0][0]
dropout_19 (Dropout)	(None, 512)	0	dense_26[0][0]
dense_27 (Dense)	(None, 256)	131328	dropout_19[0][0]
dropout_20 (Dropout)	(None, 256)	0	dense_27[0][0]
dense_28 (Dense)	(None, 128)	32896	dropout_20[0][0]
dropout_21 (Dropout)	(None, 128)	0	dense_28[0][0]
dense_29 (Dense)	(None, 2)	258	dropout_21[0][0]
=====			
Total params: 27,464,152			
Trainable params: 13,712,752			
Non-trainable params: 13,751,400			
=====			

In [76]:

```
plot_model(model, to_file='model_1.png', show_shapes=True, show_layer_names=True)
```

Out[76]:



In [71]:

```
history = model.fit(x=[X1_tr,X2_tr,X3_tr,X4_tr, X5_tr, X6_tr,X7_tr], y=y_train, validation_data=([X1_cv,X2_cv,X3_cv,X4_cv, X5_cv, X6_cv,X7_cv]),y_cv),epochs=10,batch_size=500,verbose=2)
```

Train on 64000 samples, validate on 16000 samples

```
Epoch 1/10
- 65s - loss: 0.3529 - auroc: 0.7886 - val_loss: 0.3773 - val_auroc: 0.7444
Epoch 2/10
- 64s - loss: 0.3517 - auroc: 0.7915 - val_loss: 0.3779 - val_auroc: 0.7443
Epoch 3/10
- 65s - loss: 0.3528 - auroc: 0.7879 - val_loss: 0.3777 - val_auroc: 0.7443
Epoch 4/10
- 65s - loss: 0.3505 - auroc: 0.7917 - val_loss: 0.3778 - val_auroc: 0.7446
Epoch 5/10
- 64s - loss: 0.3519 - auroc: 0.7894 - val_loss: 0.3775 - val_auroc: 0.7449
Epoch 6/10
- 65s - loss: 0.3518 - auroc: 0.7903 - val_loss: 0.3789 - val_auroc: 0.7447
Epoch 7/10
- 64s - loss: 0.3501 - auroc: 0.7939 - val_loss: 0.3789 - val_auroc: 0.7448
Epoch 8/10
- 64s - loss: 0.3506 - auroc: 0.7926 - val_loss: 0.3777 - val_auroc: 0.7451
Epoch 9/10
- 64s - loss: 0.3504 - auroc: 0.7926 - val_loss: 0.3773 - val_auroc: 0.7451
Epoch 10/10
- 65s - loss: 0.3492 - auroc: 0.7947 - val_loss: 0.3788 - val_auroc: 0.7450
```

In [0]:

```
score = model.evaluate(x=[X1 test,X2 test,X3 test,X4 test, X5 test, X6 test, X7 test], y=y test, ve
```

```
rbose=2,batch_size=500)
```

In [73]:

```
print("Test Loss:", score[0])
print("Test AUC:", score[1])
```

Test Loss: 0.3780378520488739
Test AUC: 0.7460865303064443



In [74]:

```
plt.plot(history.history['auroc'])
plt.plot(history.history['val_auroc'])

plt.title('model auc')
plt.ylabel('auc')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])

plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



In [0]:


```
# serialize weights to HDF5
model.save_weights("model_1.h5")
print("Saved model to disk")
```

Saved model to disk

Model 2

In [75]:

```
from sklearn.feature_extraction.text import TfidfVectorizer
import seaborn as sns
vectorizer = TfidfVectorizer()
vectorizer.fit(X_train['essay'].values)
plt.boxplot(list(vectorizer.idf_))
plt.xlabel('Essay')
plt.ylabel('IDF Value')
plt.show()
```



In [77]:

```
tenth_percentile=np.quantile((vectorizer.idf_),0.10)
ninty_percentile=np.quantile((vectorizer.idf_),0.90)
print(tenth_percentile)
print(ninty_percentile)
dictionary = dict(zip(vectorizer.get_feature_names(), list(vectorizer.idf_)))
filtered_words=[]
for k,v in dictionary.items():
    if v > 2.0 and v < 11.0:
        filtered_words.append(k)

len(filtered_words)
```

```
7.375306104990596
11.373506806659794
```

Out[77]:

```
27930
```

In [0]:

```
tokenizer = Tokenizer()
tokenizer.fit_on_texts(filtered_words)

X8_tr = np.array(tokenizer.texts_to_sequences(X_train['essay'].values))
X8_cv = np.array(tokenizer.texts_to_sequences(X_cv['essay'].values))
X8_test = np.array(tokenizer.texts_to_sequences(X_test['essay'].values))
```

In [0]:

```
vocab_size = len(tokenizer.word_index) + 1

maxlen = 200

X8_tr = pad_sequences(X8_tr, padding='post', maxlen=maxlen)
X8_cv = pad_sequences(X8_cv, padding='post', maxlen=maxlen)
X8_test = pad_sequences(X8_test, padding='post', maxlen=maxlen)
```

In [80]:

```
print(X8_tr.shape)
print(X8_cv.shape)
print(X8_test.shape)
```

```
(64000, 200)
(16000, 200)
(20000, 200)
```

In [0]:

```
with open('glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())
```

In [0]:

```
embeddings_dictionary = dict()
for word in glove_words:
    vector_dimensions = model[word]
    embeddings_dictionary[word] = vector_dimensions
```

In [0]:

```
embedding_matrix = np.zeros((vocab_size, 300))
for word, index in tokenizer.word_index.items():
    embedding_vector = embeddings_dictionary.get(word)
    if embedding_vector is not None:
        embedding_matrix[index] = embedding_vector
```

In [84]:

```
embedding_matrix.shape
```

Out[84]:

```
(27931, 300)
```

In [85]:

```
input_8 = Input(shape=(maxlen,), name='essay_tfidf_input')
print(input_8.shape)
input_8_embedding = Embedding(vocab_size, 300, weights=[embedding_matrix], trainable=False)(input_8)
print(input_8_embedding.shape)
input_8_lstm = LSTM(128, return_sequences=True)(input_8_embedding)
print(input_8_lstm.shape)
input_8_flatten = Flatten()(input_8_lstm)
print(input_8_flatten.shape)
```

```
(?, 200)
(?, 200, 300)
(?, ?, 128)
(?, ?)
```

In [0]:

```
#At the end we concatenate altogether and add other Dense layers
output_2 = Concatenate()
([input_8_flatten, input_2_flatten, input_3_flatten, input_4_flatten, input_5_flatten, input_6_flatten,
input_7_dense])
output_2 = BatchNormalization()(output_2)
output_2 = Dense(512, kernel_initializer="he_uniform", activation='relu')(output_2)
output_2 = BatchNormalization()(output_2)
output_2 = Dropout(0.5)(output_2)
output_2 = Dense(256, kernel_initializer="he_uniform", activation='relu')(output_2)
output_2 = BatchNormalization()(output_2)
```

```

output_2 = BatchNormalization()(output_2)
output_2 = Dropout(0.5)(output_2)
output_2 = Dense(128, kernel_initializer="he_uniform", activation='relu')(output_2)
output_2 = BatchNormalization()(output_2)
output_2 = Dropout(0.5)(output_2)
output_2 = Dense(2, activation='softmax')(output_2)

```

In [95]:

```

model_2 = Model(inputs=[input_8,input_2,input_3,input_4,input_5,input_6,input_7], outputs=output_2)
model_2.compile(loss='binary_crossentropy', optimizer=optimizers.SGD(lr=0.001), metrics=[auroc])
model_2.summary()

```

Model: "model_13"

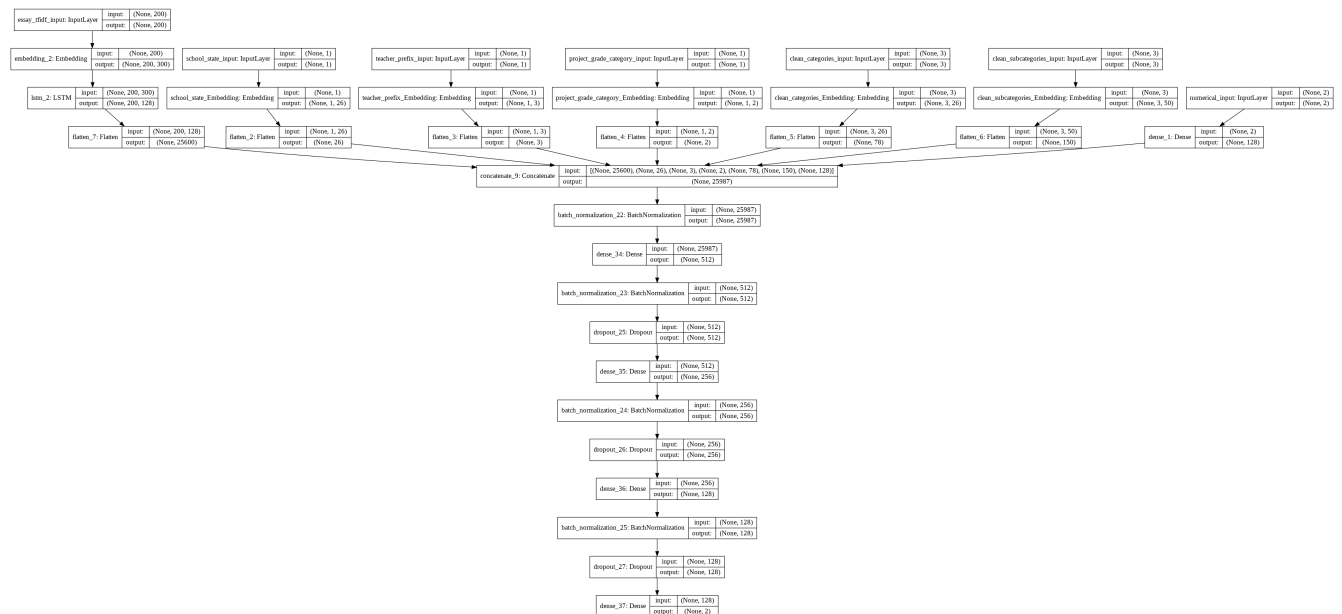
Layer (type)	Output Shape	Param #	Connected to
essay_tfidf_input (InputLayer)	(None, 200)	0	
embedding_2 (Embedding)	(None, 200, 300)	8379300	essay_tfidf_input[0][0]
school_state_input (InputLayer)	(None, 1)	0	
teacher_prefix_input (InputLayer)	(None, 1)	0	
project_grade_category_input (InputLayer)	(None, 1)	0	
clean_categories_input (InputLayer)	(None, 3)	0	
clean_subcategories_input (InputLayer)	(None, 3)	0	
lstm_2 (LSTM)	(None, 200, 128)	219648	embedding_2[0][0]
school_state_Embedding (Embedding)	(None, 1, 26)	1352	school_state_input[0][0]
teacher_prefix_Embedding (Embedding)	(None, 1, 3)	18	teacher_prefix_input[0][0]
project_grade_category_Embedding (Embedding)	(None, 1, 2)	10	project_grade_category_input[0][0]
clean_categories_Embedding (Embedding)	(None, 3, 26)	1352	clean_categories_input[0][0]
clean_subcategories_Embedding (Embedding)	(None, 3, 50)	19650	clean_subcategories_input[0][0]
numerical_input (InputLayer)	(None, 2)	0	
flatten_7 (Flatten)	(None, 25600)	0	lstm_2[0][0]
flatten_2 (Flatten)	(None, 26)	0	school_state_Embedding[0][0]
flatten_3 (Flatten)	(None, 3)	0	teacher_prefix_Embedding[0][0]
flatten_4 (Flatten)	(None, 2)	0	project_grade_category_Embedding[0][0]
flatten_5 (Flatten)	(None, 78)	0	clean_categories_Embedding[0][0]
flatten_6 (Flatten)	(None, 150)	0	clean_subcategories_Embedding[0][0]
dense_1 (Dense)	(None, 128)	384	numerical_input[0][0]
concatenate_9 (Concatenate)	(None, 25987)	0	flatten_7[0][0] flatten_2[0][0] flatten_3[0][0] flatten_4[0][0] flatten_5[0][0] flatten_6[0][0] dense_1[0][0]
batch_normalization_22 (Batch Normalization)	(None, 25987)	103948	concatenate_9[0][0]
dense_34 (Dense)	(None, 512)	13305856	batch_normalization_22[0][0]
batch_normalization_23 (Batch Normalization)	(None, 512)	2048	dense_34[0][0]
dropout_25 (Dropout)	(None, 512)	0	batch_normalization_23[0][0]

dense_35 (Dense)	(None, 256)	131328	dropout_25[0][0]
batch_normalization_24 (Batch Normalization)	(None, 256)	1024	dense_35[0][0]
dropout_26 (Dropout)	(None, 256)	0	batch_normalization_24[0][0]
dense_36 (Dense)	(None, 128)	32896	dropout_26[0][0]
batch_normalization_25 (Batch Normalization)	(None, 128)	512	dense_36[0][0]
dropout_27 (Dropout)	(None, 128)	0	batch_normalization_25[0][0]
dense_37 (Dense)	(None, 2)	258	dropout_27[0][0]
=====			
Total params: 22,199,584			
Trainable params: 13,766,518			
Non-trainable params: 8,433,066			

In [97]:

```
plot_model(model_2, to_file='model_2.png', show_shapes=True, show_layer_names=True)
```

Out[97]:



In [102]:

```
history = model_2.fit(x=[X8_tr,X2_tr,X3_tr,X4_tr, X5_tr, X6_tr,X7_tr], y=y_train, validation_data=(
[X8_cv,X2_cv,X3_cv,X4_cv, X5_cv, X6_cv,X7_cv],y_cv),epochs=30,batch_size=500,verbose=2)
```

Train on 64000 samples, validate on 16000 samples

```
Epoch 1/30
- 67s - loss: 0.4107 - auroc: 0.6641 - val_loss: 0.4014 - val_auroc: 0.6906
Epoch 2/30
- 68s - loss: 0.4110 - auroc: 0.6649 - val_loss: 0.4012 - val_auroc: 0.6908
Epoch 3/30
- 67s - loss: 0.4113 - auroc: 0.6662 - val_loss: 0.4011 - val_auroc: 0.6911
Epoch 4/30
- 67s - loss: 0.4104 - auroc: 0.6671 - val_loss: 0.4007 - val_auroc: 0.6916
Epoch 5/30
- 67s - loss: 0.4098 - auroc: 0.6655 - val_loss: 0.4003 - val_auroc: 0.6920
Epoch 6/30
- 67s - loss: 0.4099 - auroc: 0.6655 - val_loss: 0.4000 - val_auroc: 0.6924
Epoch 7/30
- 67s - loss: 0.4094 - auroc: 0.6701 - val_loss: 0.3998 - val_auroc: 0.6927
Epoch 8/30
- 67s - loss: 0.4076 - auroc: 0.6718 - val_loss: 0.3995 - val_auroc: 0.6928
Epoch 9/30
- 67s - loss: 0.4078 - auroc: 0.6731 - val_loss: 0.3996 - val_auroc: 0.6930
Epoch 10/30
- 67s - loss: 0.4085 - auroc: 0.6694 - val_loss: 0.3992 - val_auroc: 0.6931
```

```

Epoch 11/30
- 67s - loss: 0.4078 - auroc: 0.6706 - val_loss: 0.3992 - val_auroc: 0.6934
Epoch 12/30
- 67s - loss: 0.4053 - auroc: 0.6772 - val_loss: 0.3989 - val_auroc: 0.6937
Epoch 13/30
- 67s - loss: 0.4069 - auroc: 0.6733 - val_loss: 0.3988 - val_auroc: 0.6940
Epoch 14/30
- 67s - loss: 0.4071 - auroc: 0.6723 - val_loss: 0.3984 - val_auroc: 0.6942
Epoch 15/30
- 67s - loss: 0.4047 - auroc: 0.6793 - val_loss: 0.3983 - val_auroc: 0.6945
Epoch 16/30
- 67s - loss: 0.4056 - auroc: 0.6762 - val_loss: 0.3981 - val_auroc: 0.6949
Epoch 17/30
- 67s - loss: 0.4062 - auroc: 0.6740 - val_loss: 0.3981 - val_auroc: 0.6952
Epoch 18/30
- 67s - loss: 0.4057 - auroc: 0.6758 - val_loss: 0.3979 - val_auroc: 0.6953
Epoch 19/30
- 68s - loss: 0.4042 - auroc: 0.6798 - val_loss: 0.3978 - val_auroc: 0.6955
Epoch 20/30
- 67s - loss: 0.4042 - auroc: 0.6794 - val_loss: 0.3976 - val_auroc: 0.6957
Epoch 21/30
- 67s - loss: 0.4034 - auroc: 0.6799 - val_loss: 0.3973 - val_auroc: 0.6960
Epoch 22/30
- 67s - loss: 0.4044 - auroc: 0.6786 - val_loss: 0.3972 - val_auroc: 0.6962
Epoch 23/30
- 67s - loss: 0.4040 - auroc: 0.6784 - val_loss: 0.3971 - val_auroc: 0.6964
Epoch 24/30
- 68s - loss: 0.4023 - auroc: 0.6848 - val_loss: 0.3970 - val_auroc: 0.6966
Epoch 25/30
- 67s - loss: 0.4031 - auroc: 0.6815 - val_loss: 0.3967 - val_auroc: 0.6969
Epoch 26/30
- 67s - loss: 0.4029 - auroc: 0.6819 - val_loss: 0.3963 - val_auroc: 0.6971
Epoch 27/30
- 67s - loss: 0.4021 - auroc: 0.6843 - val_loss: 0.3960 - val_auroc: 0.6974
Epoch 28/30
- 68s - loss: 0.4026 - auroc: 0.6824 - val_loss: 0.3962 - val_auroc: 0.6975
Epoch 29/30
- 67s - loss: 0.4023 - auroc: 0.6829 - val_loss: 0.3961 - val_auroc: 0.6977
Epoch 30/30
- 67s - loss: 0.4012 - auroc: 0.6866 - val_loss: 0.3959 - val_auroc: 0.6981

```

In [0]:

```

score = model_2.evaluate(x=[X8_test,X2_test,X3_test,X4_test, X5_test, X6_test, X7_test], y=y_test,
verbose=2,batch_size=500)

```

In [104]:

```

print("Test Loss:", score[0])
print("Test AUC:", score[1])

```

```

Test Loss: 0.39928831085562705
Test AUC: 0.6863849997066904

```

In [105]:

```



plt.plot(history.history['auroc'])
plt.plot(history.history['val_auroc'])

plt.title('model auc')
plt.ylabel('auc')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])

plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

```

In [0]:

```
# serialize weights to HDF5
model_2.save_weights("model_2.h5")
print("Saved model to disk")
```

Saved model to disk

Model 3

In [106]:

```
numeric_tr = np.hstack((X2_tr,X3_tr,X4_tr,X5_tr,X6_tr,X7_tr))
numeric_cv = np.hstack((X2_cv,X3_cv,X4_cv,X5_cv,X6_cv,X7_cv))
numeric_test = np.hstack((X2_test,X3_test,X4_test,X5_test,X6_test,X7_test))
print(numeric_tr.shape)
print(numeric_cv.shape)
print(numeric_test.shape)
```

```
(64000, 11)
(16000, 11)
(20000, 11)
```

In [107]:

```
numeric_tr=np.expand_dims(numeric_tr,axis=2)
numeric_cv=np.expand_dims(numeric_cv,axis=2)
numeric_test=np.expand_dims(numeric_test,axis=2)
print(numeric_tr.shape)
print(numeric_cv.shape)
print(numeric_test.shape)
```

```
(64000, 11, 1)
(16000, 11, 1)
(20000, 11, 1)
```

```
(10000, 11, 1),  
(20000, 11, 1))
```

In [108]:

```
input_9 = Input(shape=(numeric_tr.shape[1],numeric_tr.shape[2],),name='combined_input')  
print(input_9.shape)
```

```
(?, 11, 1)
```

In [109]:

```
#At the end we concatenate altogether and add other Dense layers  
  
#output_3=tf.reshape(output_3,[-1,output_3.shape[1],output_3.shape[1]])  
#print(output_3.shape)  
output_3 = Conv1D(128, 3, strides=1,activation="relu",kernel_initializer='he_uniform',padding='same') (input_9)  
output_3 = MaxPooling1D(pool_size=2) (output_3)  
output_3= Dropout(0.2) (output_3)  
output_3 = Conv1D(64, 3, activation="relu") (output_3)  
output_3 = MaxPooling1D(pool_size=2) (output_3)  
output_3= Dropout(0.4) (output_3)  
output_3 = Flatten() (output_3)  
  
output_4 = Concatenate() ([input_1_flatten,output_3])  
output_4 = Dense(256, kernel_initializer="he_uniform",activation='relu') (output_4)  
output_4= Dropout(0.2) (output_4)  
output_4 = Dense(128, kernel_initializer="he_uniform",activation='relu') (output_4)  
output_4= Dropout(0.3) (output_4)  
output_4 = Dense(64, kernel_initializer="he_uniform", activation='relu') (output_4)  
output_4= Dropout(0.4) (output_4)  
output_4 = Dense(2, activation='softmax') (output_4)
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4267: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

In [110]:

```
from keras import optimizers  
model_3 = Model(inputs=[input_1,input_2,input_3,input_4,input_5,input_6,input_9], outputs=output_4)  
model_3.compile(loss='binary_crossentropy', optimizer=optimizers.SGD(lr=0.001, decay=1e-6, momentum=0.9),metrics=[auroc])  
model_3.summary()
```

Model: "model_14"

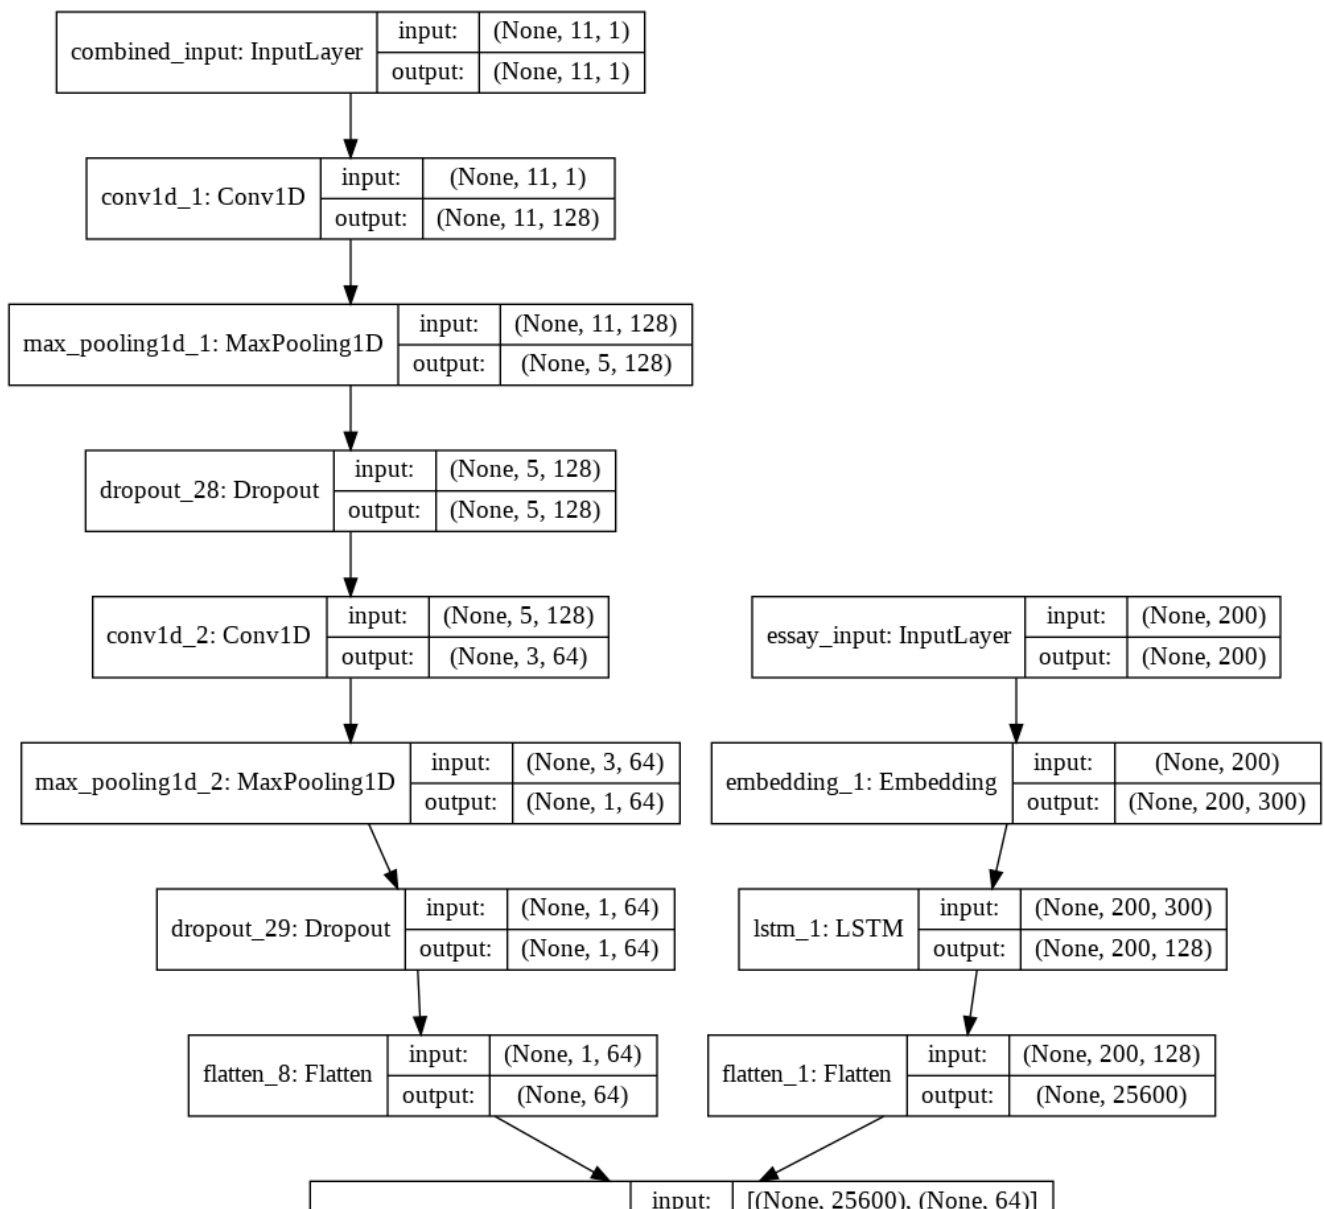
Layer (type)	Output Shape	Param #	Connected to
=====			
combined_input (InputLayer)	(None, 11, 1)	0	
conv1d_1 (Conv1D)	(None, 11, 128)	512	combined_input[0][0]
max_pooling1d_1 (MaxPooling1D)	(None, 5, 128)	0	conv1d_1[0][0]
dropout_28 (Dropout)	(None, 5, 128)	0	max_pooling1d_1[0][0]
essay_input (InputLayer)	(None, 200)	0	
conv1d_2 (Conv1D)	(None, 3, 64)	24640	dropout_28[0][0]
embedding_1 (Embedding)	(None, 200, 300)	13751400	essay_input[0][0]
max_pooling1d_2 (MaxPooling1D)	(None, 1, 64)	0	conv1d_2[0][0]
lstm_1 (LSTM)	(None, 200, 128)	219648	embedding_1[0][0]
dropout_29 (Dropout)	(None, 1, 64)	0	max_pooling1d_2[0][0]
flatten_1 (Flatten)	(None, 25600)	0	lstm_1[0][0]

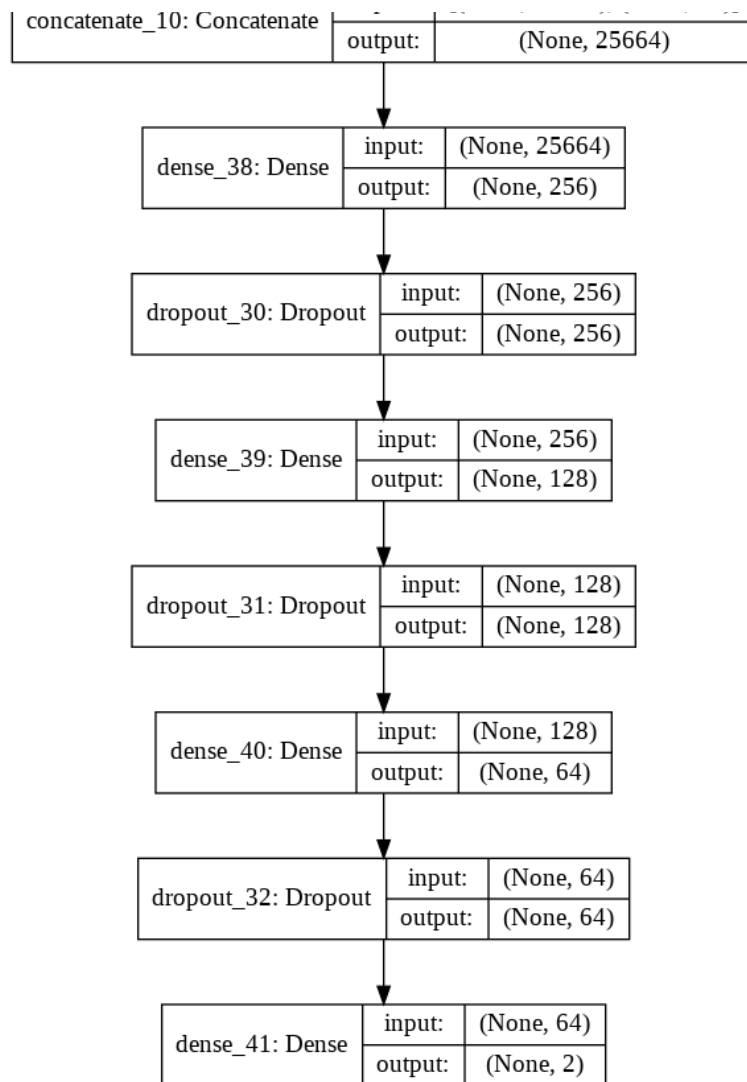
flatten_1 (Flatten)	(None, 25600)	0	flatten_1[0][0]
flatten_8 (Flatten)	(None, 64)	0	dropout_29[0][0]
concatenate_10 (Concatenate)	(None, 25664)	0	flatten_1[0][0] flatten_8[0][0]
dense_38 (Dense)	(None, 256)	6570240	concatenate_10[0][0]
dropout_30 (Dropout)	(None, 256)	0	dense_38[0][0]
dense_39 (Dense)	(None, 128)	32896	dropout_30[0][0]
dropout_31 (Dropout)	(None, 128)	0	dense_39[0][0]
dense_40 (Dense)	(None, 64)	8256	dropout_31[0][0]
dropout_32 (Dropout)	(None, 64)	0	dense_40[0][0]
dense_41 (Dense)	(None, 2)	130	dropout_32[0][0]
=====			
Total params: 20,607,722			
Trainable params: 6,856,322			
Non-trainable params: 13,751,400			

In [111]:

```
plot_model(model_3, to_file='model_3.png', show_shapes=True, show_layer_names=True)
```

Out[111]:





In [0]:

```
history = model_3.fit(x=[X1_tr,X2_tr,X3_tr,X4_tr, X5_tr, X6_tr,numeric_tr], y=y_train, validation_data=([X1_cv,X2_cv,X3_cv,X4_cv, X5_cv, X6_cv,numeric_cv],y_cv),epochs=15,batch_size=300,verbose=2)
```

Train on 64000 samples, validate on 16000 samples

Epoch 1/15

- 88s - loss: 0.3861 - auroc: 0.7274 - val_loss: 0.3936 - val_auroc: 0.7068

Epoch 2/15

- 87s - loss: 0.3849 - auroc: 0.7293 - val_loss: 0.3943 - val_auroc: 0.7080

Epoch 3/15

- 87s - loss: 0.3842 - auroc: 0.7312 - val_loss: 0.4000 - val_auroc: 0.7085

Epoch 4/15

- 88s - loss: 0.3826 - auroc: 0.7349 - val_loss: 0.3901 - val_auroc: 0.7096

Epoch 5/15

- 87s - loss: 0.3820 - auroc: 0.7371 - val_loss: 0.3935 - val_auroc: 0.7106

Epoch 6/15

- 87s - loss: 0.3795 - auroc: 0.7418 - val_loss: 0.3895 - val_auroc: 0.7110

Epoch 7/15

- 88s - loss: 0.3791 - auroc: 0.7444 - val_loss: 0.3893 - val_auroc: 0.7106

Epoch 8/15

- 87s - loss: 0.3778 - auroc: 0.7474 - val_loss: 0.3908 - val_auroc: 0.7110

Epoch 9/15

- 87s - loss: 0.3761 - auroc: 0.7490 - val_loss: 0.3902 - val_auroc: 0.7121

Epoch 10/15

- 86s - loss: 0.3743 - auroc: 0.7549 - val_loss: 0.3898 - val_auroc: 0.7119

Epoch 11/15

- 87s - loss: 0.3729 - auroc: 0.7576 - val_loss: 0.3910 - val_auroc: 0.7098

Epoch 12/15

- 87s - loss: 0.3706 - auroc: 0.7621 - val_loss: 0.3947 - val_auroc: 0.7110

Epoch 13/15

- 87s - loss: 0.3710 - auroc: 0.7638 - val_loss: 0.3911 - val_auroc: 0.7091

Epoch 14/15

- 87s - loss: 0.3682 - auroc: 0.7679 - val_loss: 0.3907 - val_auroc: 0.7090

Epoch 15/15

```
Epoch 10/10  
- 87s - loss: 0.3657 - auroc: 0.7726 - val_loss: 0.3914 - val_auroc: 0.7092
```

In [0]:

```
score = model_3.evaluate(x=[X1_test,X2_test,X3_test,X4_test, X5_test, X6_test, numeric_test], y=y_test, verbose=2,batch_size=300)
```



In [0]:

```
print("Test Loss:", score[0])  
print("Test AUC:", score[1])
```

```
Test Loss: 0.39427627846598623  
Test AUC: 0.7000997626670504
```

In [0]:

```
plt.plot(history.history['auroc'])  
plt.plot(history.history['val_auroc'])  
  
plt.title('model auc')  
plt.ylabel('auc')  
plt.xlabel('epoch')  
plt.legend(['train', 'test'], loc='upper left')  
plt.show()  
  
plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])  
  
plt.title('model loss')  
plt.ylabel('loss')  
plt.xlabel('epoch')  
plt.legend(['train', 'test'], loc='upper left')  
plt.show()
```



In [0]:

```
# serialize weights to HDF5
model_3.save_weights("model_3.h5")
print("Saved model to disk")
```

Saved model to disk